

Metacasanova

A High-performance Meta-compiler for
Domain-specific Languages



Francesco Di Giacomo

Metacasanova: a High-performance Meta-compiler for Domain-specific Languages

Francesco Di Giacomo

Contents

1	Introduction	1
1.1	Algorithms and problems	2
1.2	Programming languages	3
1.2.1	Low-level programming languages	3
1.2.2	High-level programming languages	5
1.2.3	General-purpose vs Domain-specific languages	7
1.3	Compilers	9
1.4	Meta-compilers	10
1.4.1	Requirements	11
1.4.2	Benefits	11
1.4.3	Scientific relevance	12
1.5	Problem statement	13
1.6	Thesis structure	14
2	Background	15
2.1	Architectural overview of a compiler	15
2.2	Lexer	16
2.2.1	Finite state automata for regular expressions	18
2.2.2	Conversion of a NFA into a DFA	18
2.3	Parser	19
2.3.1	LR(k) parsers	21
2.3.2	Parser generators	24
2.3.3	Monadic parsers	27
2.4	Type systems and type checking	30
2.5	Semantics and code generation	31
2.6	Metaprogramming and metacompilers	33
2.6.1	Template metaprogramming	33
2.6.2	Metacompilers	35
3	Metacasanova	39
3.1	Repetitive steps in compiler development	39
3.1.1	Hard-coded implementation of type rules	40
3.1.2	Hard-coded implementation of Semantics	45
3.1.3	Discussion	49
3.2	Metacasanova overview	49
3.2.1	Requirements of Metacasanova	49
3.2.2	Program structure	50

3.2.3	Formalization	51
3.3	Architectural overview	53
3.4	Parsing	54
3.4.1	Declarations	54
3.4.2	Rules	56
3.4.3	Parser post-processor	57
3.5	Type checking	63
3.5.1	Checking declarations	63
3.5.2	Checking rules	66
3.6	Code generation	69
3.6.1	Meta-data structures code generation	69
3.6.2	Code generation for rules	70
4	Language design in Metacasanova	75
4.1	The C-- language	75
4.2	Casanova 2.5 in Metacasanova	75
4.2.1	The Casanova language	75
4.2.2	Casanova 2.5	76
4.2.3	Chosen languages	79
4.2.4	Performance	79
4.2.5	Discussion	79
5	Metacasanova optimization	83
5.1	Case study	83
5.2	Using Modules and Functors in Metacasanova	84
5.3	Functor result inlining	86
5.4	Functor interpreter	88
5.5	C-- optimization	89
5.6	Casanova 2.5 optimization	89
5.7	Evaluation	89
6	Networking primitives in Casanova 2	93
6.1	Introduction	93
6.2	Motivation	95
6.3	Related work	95
6.4	The master/slave network architecture	96
6.5	Case study	98
6.6	Implementation	99
6.7	Networking in Metacasanova	103
7	Discussion and conclusion	105
A	List operations with templates	107

Chapter 1

Introduction

About the use of language: it is impossible to sharpen a pencil with a blunt axe. It is equally vain to try to do it with ten blunt axes instead.

Edsger Dijkstra

The number of programming languages available on the market has dramatically increased during the last years. The tiobe index [...], a ranking of programming languages based on their popularity, lists 50 programming languages for 2017. This number is only a small glimpse of the real amount, since it does not take into account several languages dedicated to specific applications. This growth has brought a further need for new compilers that are able to translate programs written in those languages into executable code. The goal of this work is to investigate how the development speed of a compiler can be boosted by employing meta-compilers, programs that generalize the task performed by a normal compiler. In particular the goal of this research is creating a meta-compiler that significantly reduces the amount of code needed to define a language and its compilation steps, while maintaining acceptable performance.

This chapter introduces the issue of expressing the solution of problems in terms of algorithms in Section 1.1. Then we proceed by defining how the semi-formal definition of an algorithm must be translated into code executable by a processor (Section 1.2). In this section we discuss the advantages and disadvantages of using different kinds of programming languages with respect to their affinity with the specific hardware architecture and the scope of the domain they target. In Section 1.3 we explain the reason behind compilers and we explain why building a compiler is a time-consuming task. In Section 1.4 we introduce the idea of meta-compilers as a further step into generalizing the task of compilers. In this section we also explain the requirements, benefits, and the relevance as a scientific topic. Finally in Section 1.5 we formulate the problem statement and the research questions that this work will answer.




Figure 1.1: Flow chart for the sum of a sequence of numbers

1.1 Algorithms and problems

Since the ancient age, there has always been the need of describing the sequence of activities needed to perform a specific task [...], to which we refer with the name of *Algorithm*. The most ancient known example of this dates back to the Babylonians, who invented algorithms to perform the factorization and the approximation of the square root [...]. Regardless of the specific details of each algorithm, one needs to use some kind of language to define the sequence of steps to perform. In the past people used natural language to describe such steps but, with the advent of the computer era, the choice of the language has been strictly connected with the possibility of its implementation [...]. Natural languages are not suitable for the implementation, as they are known to be verbose and ambiguous. For this reason, several kind of formal solutions have been employed, which are described below.

Flow charts

A flow chart is a diagram where the steps of an algorithm are defined by using boxes of different kinds, connected by arrows to define their ordering in the sequence. The boxes are rectangular-shaped if they define an *activity* (or processing step), while they are diamond-shaped if they define a *decision*. A rectangle with rounded corners denotes the initial step. An example of a flow chart describing how to sum the numbers in a sequence is described in Figure 1.1.

Pseudocode

Pseudocode is a semi-formal language that might contain also statements expressed in natural language and omits system specific code like opening file writers, printing messages on the standard output, or even some data

structure declaration and initialization. It is intended mainly for human reading rather than machine reading. The pseudocode to sum a sequence of numbers is shown in Algorithm 1.1.

Algorithm 1.1 Pseudocode to perform the sum of a sequence of integer numbers

```

function SUMINTEGERS(l list of integers)
    sum  $\leftarrow$  0
    for all x in l do
        sum  $\leftarrow$  sum + x
    end for
    return sum
end function

```

Advantages and disadvantages

Using flow charts or pseudo-code has the advantage of being able to define an algorithm in a way which is very close to the abstractions employed when using natural language: a flow chart combines both the use of natural language and a visual interface to describe an algorithm, pseudo-code allows to employ several abstractions and even define some steps in terms of natural language. The drawback of these two formal representations is that, when it comes to the implementation, the definition of the algorithm must be translated by hand into code that the hardware is able to execute. This could be done by implementing the algorithm in a low-level or high-level programming language. This process affects at different levels how the logic of the algorithm is presented, as explained further.

1.2 Programming languages

A programming language is a formal language that is used to define instructions that a machine, usually a computer, must perform in order to produce a result through computation [...]. There is a variety of taxonomies used to classify programming languages [...], but all of them are considering two main characteristics [...]: the level of abstraction, or how close to the specific targeted hardware they are, and the domain, which defines the range of applicability of a programming language. In the following sections we give an exhaustive explanation of the aforementioned characteristics.

1.2.1 Low-level programming languages

A low-level programming language is a programming language that provides little to no abstraction from the hardware architecture of a processor [...]. This means that it is strongly connected with the instruction set of the targeted machine, the set of instructions a processor is able to execute. These languages are divided into two sub-categories: *first-generation* and *second-generation* languages [...]:



Figure 1.2: Front panel of IBM 1620

First-generation languages

Machine code falls into the category of first-generation languages. In this category we find all those languages that do not require code transformations to be executed by the processor. These languages were used mainly during the dawn of computer age and are rarely employed by programmers nowadays [...]. Machine code is made of stream of binary data, that represents the instruction codes and their arguments [...]. Usually this stream of data is treated by programmers in hexadecimal format, which is then remapped into binary code. The programs written in machine code were once loaded into the processor through a front panel, a controller that allowed the display and alteration of the registers and memory (see Figure 1.2). An example of machine code for a program that computes the sum of a sequence of integer numbers can be seen in Listing 1.1.

```

1  00075 c7 45 b8 00 00
2  00 00
3  0007c eb 09
4  0007e 8b 45 b8
5  00081 83 c0 01
6  00084 89 45 b8
7  00087 83 7d b8 0a
8  0008b 7d 0f
9  0008d 8b 45 b8
10 00090 8b 4d c4
11 00093 03 4c 85 d0
12 00097 89 4d c4
13 0009a eb e2

```

Listing 1.1: Machine code to compute the sum of a sequence of numbers

Second-generation languages

The languages in this category provides an abstraction layer over the machine code by expressing processor instructions with mnemonic names both for the instruction code and the arguments. For example, the arithmetic sum instruction `add` is the mnemonic name for the instruction code `0x00` in x86 processors. Among these languages we find *Assembly*, that is mapped

with an *Assembler* to machine code. The Assembler can load directly the code or link different *object files* to generate a single executable by using a *linker*. An example of assembly x86 code corresponding to the machine code in Listing 1.1 can be found in Listing 1.2. You can see that the code in the machine code 00081 83 c0 01 at line 5 has been replaced by its mnemonic representation in Assembly as `add eax, 1`.

```

1  mov DWORD PTR _i$1[ebp], 0
2  jmp SHORT $LN4@main
3  $LN2@main:
4  mov eax, DWORD PTR _i$1[ebp]
5  add eax, 1
6  mov DWORD PTR _i$1[ebp], eax
7  $LN4@main:
8  cmp DWORD PTR _i$1[ebp], 10      ; 0000000aH
9  jge SHORT $LN3@main
10 mov eax, DWORD PTR _i$1[ebp]
11 mov ecx, DWORD PTR _sum$[ebp]
12 add ecx, DWORD PTR _numbers$[ebp+eax*4]
13 mov DWORD PTR _sum$[ebp], ecx
14 jmp SHORT $LN2@main

```

Listing 1.2: Assembly x86 code to compute the sum of a sequence of numbers

Advantages and disadvantages

Writing a program in low-level programming languages has been known to produce programs that are generally more efficient than their high-level counterparts [...]. However, the high-performance comes at great costs: *(i)* the programmer must be an expert of the underlying architecture and of the specific instruction set of the processor, *(ii)* the program loses portability because the low-level code is tightly bound to the specific hardware architecture it targets, *(iii)* the logic and readability of the program is hidden among the details of the instruction set itself, and *(iv)* developing a program in assembly requires a considerable effort in terms of time and debugging [...]: assembly lacks any abstraction from the concrete hardware architecture, such as a type system, that partially ensures the correctness of the program or high-level constructs that allow to manipulate the execution of the program.

1.2.2 High-level programming languages

A high-level programming language is a programming language that offers a high level of abstraction from the specific hardware architecture of the machine [...]. Unlike machine code (and in some way also assembly), high-level languages are not directly executable by the processor and they require some kind of translation process into machine code. The level of abstraction offered by the language defines how high level the language is. Several categories of high-level programming language exist, but the main one are described below.

Imperative programming languages

Imperative programming languages model the computation as a sequence of statements that alter the state of the program (usually the memory state).

A program in such languages consists then of a sequence of *commands*. Notable examples are FORTRAN, C, and PASCAL. An example of the program used in Listing 1.1 and 1.2 written in C can be seen in Listing 1.3. Line 5 to 9 corresponds to the Assembly code in Listing 1.2.

```

1  int main()
2  {
3      int numbers[10] = { 1, 6, 8, -2, 4, 3, 0, 1, 10, -5 };
4      int sum = 0;
5      for (int i = 0; i < 10; i++)
6      {
7          sum += numbers[i];
8      }
9      printf("%d\n", sum);
10 }
```

Listing 1.3: C code to compute the sum of a sequence of numbers

Declarative programming languages

Declarative programming languages are antithetical to those based on imperative programming, as they model computation as an evaluation of expressions and not as a sequence of commands to execute. Declarative programming languages are called as such when they are side-effects free or referentially transparent. The definition of referential transparency varies [37], but it is usually explained with the substitution principle, which states that a language is referentially transparent if any expression can be replaced by its value without altering the behaviour of the program [34]. For instance, the following sentences in natural language are both true

```

Cicero = Tullius
''Cicero'' contains six letters
```

but they are not referentially transparent, since replacing the last name with the middle name falsifies the second sentence.

A similar situation in programming languages is met when considering variable assignments: the statement

```
x = x + 5
```

is not referentially transparent. Let us assume this statement appears twice in a program and that at the beginning $x = 0$. Clearly the expression $x + 5$ results in the value 5 the first time, but the second time the same statement is executed the expression has value 10. Thus replacing all the occurrences of $x + 5$ with 5 is wrong, which is why imperative languages are not referentially transparent. A more rigorous definition of referential transparency can be found in [41].

Declarative programming languages are often compared to imperative programming languages by stating that declarative programming defines *what* to compute and not *how* to compute it. This family of languages include *functional programming*, *logic programming*, and *database query languages*. Notable examples are F#, Haskell, Prolog, SQL, and Linq (which is a query language embedded in C#). Listing 1.4 shows the code to perform the sum of a sequence of integer numbers in F# with a recursive

function. Higher-order functions, such as `fold`, allow even to capture the same recursive pattern into a single function as shown in Listing 1.5. Both implementations are referentially transparent.

```
let rec sumList l =
  match l with
  | [] -> 0
  | x :: xs -> x + (sumList xs)
```

Listing 1.4: Recursive F# code to compute the sum of a sequence of numbers

```
let sumList l = l |> List.fold (+) 0
```

Listing 1.5: F# code to compute the sum of a sequence of numbers using higher-order functions

1.2.3 General-purpose vs Domain-specific languages

General-purpose languages are defined as languages that can be used across different application domains and lack abstractions that specifically target elements of a single domain. Example of these are languages such as C, C++, C#, and Java. Although several applications are still being developed by using general-purpose programming languages, in several contexts it is more convenient to rely on *domain-specific languages*, because they offer abstractions relative to the problem domain that are unavailable in general-purpose languages. Notable examples of the use of domain-specific languages are listed below.

Graphics programming

Rendering a scene in a 3D space is often performed by relying on dedicated hardware. Modern graphics processors rely on shaders to create various effects that are rendered in the 3D scene. Shaders are written in domain-specific languages, such as GLSL or HLSL [...], that offer abstractions to compute operations at GPU level that are often used in computer graphics, such as vertices and pixel transformations, matrix multiplications, and interpolation of textures. Listing 1.6 shows the code to implement light reflections in HLSL. At line 4 you can, for example, see the use of matrix multiplication provided as a language abstraction in HLSL.

```
1  VertexShaderOutput VertexShaderSpecularFunction(VertexShaderInput input,
2      float3 Normal : NORMAL)
3  {
4      VertexShaderOutput output;
5      float4 worldPosition = mul(input.Position, World);
6      float4 viewPosition = mul(worldPosition, View);
7      output.Position = mul(viewPosition, Projection);
8      float3 normal = normalize(mul(Normal, World));
9      output.Normal = normal;
10     output.View = normalize(float4(EyePosition, 1.0f) - worldPosition);
11     return output;
12 }
```

Listing 1.6: HLSL code to compute the light reflection

Game programming

Computer games are a field where domain-specific languages are widely employed, as they contain complex behaviours that often require special constructs to model timing event-based primitives, or to execute tasks in parallel. These behaviours cannot be modelled, for performance reasons, by using threads. Therefore, in the past, domain-specific languages which provide these abstractions have been implemented. [...]. In Listing 1.7 an example of the SQF domain-specific language for the game ArmA2 is shown. This language offers abstractions to wait for a specific amount of time, to wait for a condition, and to spawn scripts that run in parallel to the callee, that you can respectively see at lines 18, 12, and 10.

```

1  "colorCorrections" ppEffectAdjust [1, pi, 0, [0.0, 0.0, 0.0, 0.0],
    [0.05, 0.18, 0.45, 0.5], [0.5, 0.5, 0.5, 0.0]];
2  "colorCorrections" ppEffectCommit 0;
3  "colorCorrections" ppEffectEnable true;
4
5  thanatos switchMove "AmovPpneMstpSrasWrflDnon";
6  [[,(position tower) nearestObject 6540,[[ "USMC_Soldier",west]],4,true
    ,[]] execVM "patrolBuilding.sqf";
7  playMusic "Intro";
8
9  titleCut ["", "BLACK FADED", 999];
10 [] Spawn
11 {
12     waitUntil{!(isNil "BIS_fnc_init")};
13     [
14         localize "STR_TITLE_LOCATION" ,
15         localize "STR_TITLE_PERSON",
16         str(date select 1) + "." + str(date select 2) + "." + str(date
            select 0)
17     ] spawn BIS_fnc_infoText;
18     sleep 3;
19     "dynamicBlur" ppEffectEnable true;
20     "dynamicBlur" ppEffectAdjust [6];
21     "dynamicBlur" ppEffectCommit 0;
22     "dynamicBlur" ppEffectAdjust [0.0];
23     "dynamicBlur" ppEffectCommit 7;
24     titleCut ["", "BLACK IN", 5];
25 };

```

Listing 1.7: ArmA 2 scripting language

Shell scripting languages

Shell scripting languages, such as the *Unix Shell script*, are used to manipulate files or user input in different ways. They generally offer abstractions to the operating system interface in the form of dedicated commands. Listing 1.8 shows an example of a program written in Unix shell script to convert an image from JPG to PNG format. At line 3 you can see the use of the statement `echo` to display a message in the standard output.

```
1  for jpg; do
2      png="${jpg%.jpg}.png"
3      echo converting "$jpg" ...
4      if convert "$jpg" jpg.to.png ; then
5          mv jpg.to.png "$png"
6      else
7          echo 'jpg2png: error: failed output saved in "jpg.to.png".' >&2
8          exit 1
9      fi
10 done
11 echo all conversions successful
12 exit 0
```

Listing 1.8: Unix shell code

Advantages and disadvantages

High-level programming languages offer a variety of abstractions over the specific hardware the program targets. The obvious advantage of this is that the programmer does not need to be an expert of the underlying hardware architecture or instruction set. A further advantage is that the available abstractions are closer to the semi-formal description of the underlying algorithm as pseudo-code. This produces two desirable effects: (i) the readability of the program is increased as the available abstractions are closer to the natural language than the equivalent machine code, and (ii) that being able to mimic the semi-formal version of an algorithm, which is generally how the algorithm is presented and on which its correctness is proven, grants a higher degree of correctness in the specific implementation [...].

The use of a high-level programming language might, in general, not achieve the same high-performance as writing the same program with a low-level programming language, but modern code-generation optimization techniques that can achieve similar performance are known [...]. A further major issue in using high-level programming languages is that the machine cannot directly execute the code, thus the use of a compiler that translates the high-level program into machine code is necessary.

The portability of a high-level programming language depends on the architecture of the underlying compiler, thus some languages are portable and the same code can be run on different machines (for example Java), while others might require to be compiled to target a specific architecture (for example C++).

1.3 Compilers

A compiler is a program that transforms source code defined in a programming language into another computer language, which usually is object code but can also be code in a high-level programming language [...]. Writing a compiler is a necessary step to implementing a high-level programming language. Indeed, a high-level programming languages, unlike low-level ones, are not executable directly by the processor and need to be translated into machine code, as stated in Section 1.2.1 and 1.2.2.

The first complete compiler was developed by IBM for the FORTRAN language and required 18 person-years for its development [...]. This clearly shows that writing a compiler is a hard and time-consuming task.

A compiler is a complex piece of software made of several components that implement a step in the translation process. The translation process performed by a compiler involves the following steps:

1. *syntactical analysis*: In this phase the compiler checks that the program is written according to the grammar rules of the language. In this phase the compiler must be able to recognize the *syntagms* of the language (the “words”) and also check if the program conforms to the syntax rules of the language through a grammar specification.
2. *type checking*: In this phase the compiler checks that a *syntactically correct program* performs operations conform to a defined *type system*. A type system is a set of rules that assign properties called types to the constructs of a computer program [...]. The use of a type system drastically reduces the chance of having bugs in a computer program [...]. This phase can be performed at compile time (*static typing*) or the generated code could contain the code to perform the type checking at runtime (*dynamic typing*).
3. *code generation*: In this phase the compiler takes the *syntactically and type-correct program* and performs the translation step. At this point an equivalent program in a target language will be generated. The target language can be object code, another high-level programming language, or even a bytecode that can be interpreted by a virtual machine.

All the previous steps are always the same regardless of the language the compiler translates from and they are not part of the creative aspect of the language design[...]. Approaches to automating the construction of the syntactical analyser are well known in literature [...], to the point that several lexer/parser generators are available for programmers, for example all those belonging to the yacc family such as `yacc` for C/C++, `fsyacc` for F#, `cup` for Java, and `Happy` for Haskell. On the other hand, developers lack a set of tools to automate the implementation of the last two steps, namely the type checking and the code generation.

For this reason, when implementing a compiler, the formal type system definition and the operational semantics, which is tightly connected to the code generation and defines how the constructs of the language behave, must be translated into the abstractions provided by the host language in which the compiler will be implemented. Other than being a time-consuming activity itself, this causes that (i) the logic of the type system and operational semantics is lost inside the abstraction of the host-language, and (ii) it is difficult to extend the language with new features.

1.4 Meta-compilers

In Section 1.3 we described how the steps involved in designing and implementing a compiler do not require creativity and are always the same, regardless of the language the compiler is built for. The first step, namely the syntactical analysis, can be automated by using one of the several lexer/parser generators available, but the implementation of a type checker

and a code generator still relies on a manual implementation. This is where meta-compilers come into play: a meta-compiler is a program that takes the source code of another program written in a specific language and the language definition itself, and generates executable code. The language definition is written in a programming language, referred to as *meta-language*, which should provide the abstractions necessary to define the syntax, type system, and operational semantics of the language, in order to implement all the steps above.

1.4.1 Requirements

As stated in Section 1.4, a meta-compiler should provide a meta-language that is able to define the syntax, type system, and operational semantics of a programming language. In Section 1.3 we discussed how methods to automate the implementation of syntactical analyser are already known in scientific literature. For this reason, in this work, we will focus exclusively on automating the implementation of the type system and of the operational semantics. Given this focus, we formulate the following requirements:

- The meta-language should provide abstractions to define the constructs of the language. This includes the possibility of defining control structures, operators with any form of prefix or infix notation, and the priority of the constructs that is used when evaluating their behaviour. Furthermore, it must be possible to define the equivalence of language constructs. For instance, an integer constant might be considered both a value and a basic arithmetic expression.
- The meta-language must be able to mimic as close as possible the formal definition of a programming language. This will bring the following benefits: (i) Implementing the language in the meta-compiler will just involve re-writing almost one-to-one the type system or the semantics of the language with little or no change; (ii) the correctness and soundness [...] of the language formal definition will be directly reflected in the implementation of the language; indeed if a meta-program allows to mimic directly the type system and semantics of the language their correctness is transferred also in the implementation, while this might not be trivial when translating them in the abstractions of a high-level programming language; (iii) any extension of the language definition can be just added as an additional rule in the type system or the semantics.
- The meta-compiler must be able to embed libraries from external languages, so that they can be used to implement specific behaviours such as networking transmission or specific data structure usage.

1.4.2 Benefits

Programming languages usually are released with a minimal (but sufficient to be Turing-complete) set of features, and later extended in functionality in successive versions. This process tends to be slow and often significant improvements or additions are only seen years after the first release. For

example, Java was released in 1996 and lacked an important feature such as Generics until 2004, when J2SE 5.0 was released. Furthermore, Java and C++ lacked a construct, which is becoming more and more important with the years [...], such as lambda abstractions until 2016, while a similar language like C# 3.0 was released with such capability in 2008. The slow rate of change of programming languages is due to the fact that every abstraction added to the language must be reflected in all the modules of its compiler: the grammar must be extended to support new syntactical rules, the type checking of the new constructs must be added, and the appropriate code generation must be implemented. Given the complexity of compilers, this process requires a huge amount of work, and it is often obstructed by the low flexibility of the compiler as piece of software, and the need for backward compatibility [...]. Using a meta-compiler would speed up the extension of an existing language because it would require only to change on paper the type system and the operational semantics, and then add the new definitions to their counterpart written in the meta-language. This process is easier because the meta-language should mimic as close as possible their behaviour. Moreover, backward compatibility is automatically granted because an older program will simply use the extended language version to be compiled by the meta-compiler.

To this we add the fact that, in general, for the same reasons, the development of a new programming language is generally faster when using a meta-compiler. This could be beneficial to the development of a high variety of domain-specific languages. Indeed, such languages are often employed in situations where the developers have little or no resources to develop a fully-fledged hard-coded compiler by hand. For instance, it is desirable for game developers to focus on aspects that are strictly tied to the game itself, for example the development of an efficient graphics engine or to improve the game logic. At the same time they would need a domain-specific language to express some behaviours typical of games, things that could be achieved by using a meta-compiler rather than on a hand-made implementation.

1.4.3 Scientific relevance

Meta-compilers have been researched since the 1980's [...] and some implementations have been proposed [...]. In general meta-compilers perform poorly compared to hard-coded compilers because they add the additional layer of abstraction of the meta-language. Moreover, a specific implementation of a compiler opens up the possibility of implementing language-specific optimizations during the code generation phase. In general we find in scientific literature a substantial effort in developing techniques to optimize the code generation for compilers [...] but not many attempts in producing optimized meta-compilers (one notable exception being [put reference to RML here]). We argue that it could be interesting to present a novel approach into optimize the code generation of a meta-compiler, which might open new horizons to the research on code generation optimization also for normal compilers. Furthermore, the growth in need for domain-specific languages [...] requires the capability of producing compilers in a short amount of time, to which a significant contribution could be given by presenting a solution based on a meta-compiler. Finally, producing a domain-specific for a field like game development, where high performance is paramount, through

a meta-compiler could prove that they can be used to produce languages with decent performance.

1.5 Problem statement

In Section 1.2 we showed the advantages of using high-level programming languages when implementing an algorithm. Among such languages, it is sometimes desirable to employ domain-specific languages that offer abstractions relative to a specific application domain (Section 1.2.3). In Section 1.3 we described the need of a compiler for such languages, and that developing one is a time-consuming activity despite the process being, in great part, non-creative. In Section 1.4 we introduced the role of meta-compilers to speed up the process of developing a compiler and we listed the requirements and the benefits that one should have. In Section 1.4.3 we explained why we believe that meta-compilers are a relevant scientific topic if coupled with the problem of developing domain-specific languages in response to the their increasing need. We can now formulate our problem statement:

Problem statement: *To what extent does the use of a meta-compiler ease the development of a domain-specific language and what is the performance trade-off?*

The first parameter we need to evaluate in order to answer this question is the size of the code reduction needed to implement the domain-specific language. At this purpose, the following research question arises:

Research question 1: *To what extend can a meta-compiler reduce the amount of code required to create a compiler for a given programming language?*

The second parameter we need to evaluate is the eventual performance loss caused by introducing the abstraction layer provided by the meta-compiler. This leads to the following research question:

Research question 2: *How much is the performance loss introduced by the meta-compiler with respect to an implementation written in a language compiled with a traditional compiler and is this loss acceptable?*

In case of a performance loss, we need to identify the cause of this performance loss and if an improvement is possible. This leads to the following research question:

Research question 3: *What is the cause of the performance degradation when employing a meta-compiler and how can this be improved?*

1.6 Thesis structure

This thesis describes the architecture of Metacasanova, a meta-compiler whose meta-language is based on operational semantics, and a possible optimization for such meta-compiler. It also shows its capabilities by implementing a small imperative language and re-implementing the existing domain-specific language for games *Casanova 2*, extending it with abstractions to express network operations for multiplayer games.

In Chapter 2 we provide background information in order to understand the choices made for this work. The chapter presents the state of the art in designing and implementing compilers and existing research on meta-compilers.

In Chapter 3 we present the architecture of Metacasanova by extensively describing the implementation of all its modules.

In Chapter 4 we show how to use Metacasanova to implement two languages: a small imperative language, and *Casanova 2*, a language for game development. At the end of the chapter we provide an evaluation of the performance of the two languages and their implementation length with respect to existing compilers, thus answering to Research Question 1 and 2.

In Chapter 5 we discuss the performance loss of the implementation of the presented languages and we propose an extension of Metacasanova that aims to improve the performance of the generated code, thus answering Research Question 3.

In Chapter 6 we propose an extension of *Casanova 2* for multiplayer game development by providing both a hard-coded compiler solution and a meta-compiler one. We then compare the two approaches performance-wise and code-length-wise.

In Chapter 7 we discuss the result and answer the research questions.

Chapter 2

Background

Trying to outsmart a compiler
defeats much of the purpose of
using one.

Kernighan and Plauger - *The
Elements of Programming
Style*.

This chapter provides background information on compiler construction and the existing knowledge on meta-compiles. The goal of this chapter is dual: (i) it provides the reader with sufficient information to understand the implementation choices done when developing Metacasanova, and (ii) outlines the complexity of the process of designing and implementing a compiler, thus giving further motivation to this research work.

In Section 2.1 we outline the general architecture of a compiler by giving a short descriptions of all its components and how they work. In Section 2.2 we give a detailed explanation about *regular expressions* necessary to define the “words” of a language, and the lexer, showing how to implement one. In Section 2.3 we introduce the notion of *context-free grammars* and we show how to implement a parser able to process the grammatical rules of such grammar. In Section 2.3.2 we present a parser generator for the language F# that has been used for the implementation of Metacasanova. In Section 2.3.3 we show an alternative to standard parsers in functional programming languages. In Section 2.4 and 2.5 we explain how a type system and semantics of a language is expressed. In Section 2.6 we introduce the concept of metaprogramming and we show examples using metaprogramming in the abstractions provided by a general purpose language (C++), and with existing dedicated metacompilers.

2.1 Architectural overview of a compiler

Compilers are software that read as input a program written in a programming language, called *source language*, and translate it into an equivalent

program expressed with another programming language, called *target language*. Usually the target language is machine code, but this is not mandatory. A special kind of compilers are interpreted, that directly execute the program written in the source language rather than translating it into a target language. Some languages, like Java, use a hybrid approach, that is they compile the program into an intermediate language that is later interpreted by a *virtual machine*. Another approach involves the translation into a target high-level language [...].

Although the architecture of a compiler may slightly vary depending on the specific implementation, the translation process usually consists of the following steps:

1. **Lexical analysis:** this phase is performed by a module called *lexer* that is able to process the text and identify the syntactical elements of the language, called *tokens*.
2. **Syntactical analysis:** this phase is performed by a module called *parser*, that checks whether the program written in the source language is compliant to the formal syntax of the language. The parser is tightly coupled with the lexer, as it needs to identify the tokens of the language to correctly process the syntax rules. The parser outputs a representation of the program, called *Abstract Syntax Tree*, for later use.
3. **Type checking:** this phase is performed by the *type checker* that uses the rules defined by a *type system* to assign a property to the elements of the language called *type*. The types are used to determine whether the abstractions of the language, in a program that is syntactically correct, are used in a meaningful way.
4. **Code generation:** the code generation phase requires to choose one or more target languages to emit. In the latter case, the code generator must have a modular structure to allow to interchange the output language. For this reason this step is usually preceded by an *intermediate code generation* step, that converts the source program into an intermediate representation close to the target language. This phase can later be followed by different kinds of code optimization phases.

In what follows we extensively describe each module that was summarized above.

2.2 Lexer

As stated above, the lexer task is to recognize the *words* or *tokens* of the source language. In order to perform this task the token structure must be expressed in a formal way. Below we present such formalization and we describe the algorithm that actually recognizes the token.

Let us consider a finite alphabet Σ , a *language* is a set of strings, intended as sequences of characters in Σ .

Definition 2.1. A string in a language L in the alphabet Σ is a tuple of characters $\mathbf{a} \in \Sigma^n$.

A notable difference between languages in this context and human-spoken languages is that, in the former, we do not associate a meaning to the words but we are only interested to define which words are part of the language and which are not. Regular expressions are a convenient formalization to define the structure of sets of strings:

Definition 2.2. The following are the possible ways to define regular expressions:

- *Empty*: The regular expression ϵ is a language containing only the empty string.
- *Symbol*: $\forall a \in \Sigma$, a is a string containing the character a .
- *Alternation*: Given two regular expressions M and N , a string in the language of $M|N$, called alternation, is the sets of strings in the language of M or N .
- *Concatenation*: Given two regular expressions M and N , a string in the language of $M \cdot N$ is the language of strings $\alpha \cdot \beta$ such as $\alpha \in M$ and $\beta \in N$.
- *Repetition*: Given a regular expression M , its Kleene Closure M^* is formed by the concatenation of zero or more strings in the language M .

The regular expressions defined in Definition 2.2 can be combined to define tokens in a language.

Regular expressions can be processed by using a finite state automaton. Informally a finite state automaton is made of a finite set of states, an alphabet Σ of which it is able to process the symbols, and a set of symbol-labelled edges that connect two states and define how to transition from one state to another. Automata can be divided into two categories: *non-deterministic finite state automata (NFA)* and *deterministic finite state automata (DFA)*. Formally we have the following definitions:

Definition 2.3. A non-deterministic finite state automaton (NFA) is made of:

- A finite set of states S .
- An alphabet Σ of input symbols.
- A state $s_0 \in S$ that is the starting state of the automaton.
- A set of states $F \subset S$ called final or accepting states.
- A set of transitions $\mathcal{T} \subseteq S \times (\Sigma \cup \{\epsilon\}) \times S$.

Definition 2.4. A deterministic finite state automaton (DFA) is a NFA where the transition is a function, i.e.

$$\begin{aligned} \tau : S \times \Sigma &\rightarrow S \\ \tau(s_i, c) &= s_j \end{aligned}$$

and $\nexists \tau(s, c_i), \tau(s, c_j) \mid c_i = c_j \ \forall i, j$.

Informally, in NFA's there might be two transitions from the same state that can process the same symbol, while in DFA's for the same state there exists one and only one transition able to process a symbol and no transition processes the empty string. Regular expressions can be converted in NFA by using translation rules. The formalization of the algorithm can be found in [33], here we just show an informal overview for brevity.

2.2.1 Finite state automata for regular expressions

In this section we present an informal overview of the translation rules for regular expressions into NFA's, and an algorithm to convert an NFA into a DFA.

Conversion for Symbols A regular expression containing just one symbol $a \in \Sigma$ can be converted by creating a transition $\tau(s_i, a) = s_j$.

Conversion for concatenation The conversion for concatenation is recursive: the base case of the recursion is the symbol conversion. The conversion of a concatenation of n symbols $a_1 a_2, \dots, a_n$ is obtained by adding a transition from the last state of the conversion for the first $n - 1$ symbols into a new state through a transition processing the n -th symbol, $\tau(s_{n-1}, a_n) = s_n$.

Conversion for alternation The alternation $M|N$ is obtained by creating an automata with a ϵ -transition into a new state, that we call s_ϵ . From s_ϵ we recursively generate the automata for both M and N . Both automata can finally reach the same state through an ϵ -transition.

Conversion for Kleene closure The Kleene Closure M^* is obtained by initially creating an ϵ -transition into a state s_ϵ . s_ϵ can recursively transition to the automaton for M , which in turn transitions through an ϵ -transition to s_ϵ .

Conversion for M^+ The regular expression M^+ contains the concatenation of one or more strings in M . This can be translated by translating $M \cdot M^*$.

Conversion for $M?$ The regular expression $M?$ is a shortcut for $M|\epsilon$, thus it can be translated by using the conversion rule for the alternation.

2.2.2 Conversion of a NFA into a DFA

As stated in Section 2.2, a NFA might have, for the same state, a set of transitions that process the same symbol (including the empty string since ϵ -transitions are allowed). This means that a NFA must be able to guess which transition to follow when trying to process a token. This is not efficient to implement in a computer, thus it is better to use a DFA where there can be only one way of processing a symbol for a given state. An algorithm to automate such conversion exists and is presented in [7] but there exists an algorithm to directly convert regular expressions into DFA's, as shown in [6]. Below we present the algorithm to convert NFA's into DFA's.

The informal idea behind the algorithm is the following: since a DFA cannot contain ϵ -transitions or transitions from one state into another containing the same symbols, we have to construct an automaton that skips the ϵ -transitions and pre-calculates the calculation of the sets of states in advance. In order to do so, we need to be able to compute the *closure* of a

set of states. Informally the closure of a set of states S is the sates that can be reached by one of the states of S through an ϵ -transition. The formal definition is given below:

Definition 2.5. The closure $\mathcal{C}(S)$ of a set of states S is defined as

- $\mathcal{C}(S) = S \cup \left(\bigcup_{s \in T} \tau(s, \epsilon) \right)$
- if $\exists \mathcal{C}'(S) \mid \mathcal{C}(S) \subseteq \mathcal{C}'(S) \Rightarrow \mathcal{C}'(S) = \mathcal{C}(S)$.

Algorithm 2.1 Closure of S

```

 $T \leftarrow S$ 
repeat
   $T' \leftarrow T$ 
   $T \leftarrow \bigcup_{s \in T'} \tau(s, \epsilon)$ 
until  $T = T'$ 

```

Algorithm 2.1 computes the closure of a set of states. Note that the algorithm termination is granted because we are considering finite-state automata.

At this point we can build the set of all possible states reachable by consuming a specific character. We call this set *edge* of a set of states d .

Definition 2.6. Let d be a set of states, then the *edge* of d is defined as

$$\mathcal{E}(d, c) = \mathcal{C} \left(\bigcup_{s \in d} \tau(s, c) \right)$$

Now we can use the *closure* and *edge* to build the DFA from a NFA.

Algorithm 2.2 performs the conversion into a DFA but we need to adjust it in order to mark the final states of the automaton. A state d is final in the DFA if it is final if any of the states in $state[d]$ is final. In addition to marking final states, we must also keep track of what token is produced in that final state.

2.3 Parser

Regular expressions are a concise declarative way to define the lexical structure of the terms of a language, but they are insufficient to describe its syntax, i.e. how to combine tokens together to make “sentences” [8, 7]. For example, trying to define an arithmetic expression with chained sums would lead to the following (recursive) regular expression:

```
expr = "(" expr "+" expr ")" | digits
```

Now we would need to replace the regular expression with itself, thus obtaining

```
expr = "(" "(" expr "+" expr ")" | digits "+" "(" expr "+" expr ")" |
  digits ")" | digits
```

Algorithm 2.2 NFA into DFA conversion

```

states[0]  $\leftarrow \emptyset$ 
states[1]  $\leftarrow \mathcal{C}(s_1)$ 
p  $\leftarrow 1$ 
j  $\leftarrow 0$ 
while j  $\leq$  p do
  for all c  $\in \Sigma$  do
    e  $\leftarrow \mathcal{E}(\text{states}[j], c)$ 
    if  $\exists i \leq p \mid e = \text{states}[i]$  then
      trans[j, c]  $\leftarrow i$ 
    else
      p  $\leftarrow p + 1$ 
      states[p]  $\leftarrow e$ 
      trans[j, c]  $\leftarrow p$ 
    end if
  end for
  j  $\leftarrow j + 1$ 
end while

```

It is easy to see that this substitution would never end, as the regular expression keeps growing at each replacement.

A compiler uses the parser module to check the syntactical structure of a program. As we will see more in depth below, the parser is tightly coupled with the lexer, which is used by it to recognize tokens. In order to present the structure of the parser, it is first necessary to introduce *context-free grammars*.

As before we consider a language as a set of tuples of characters taken from a finite alphabet Σ . Informally, a context-free grammar is a set of productions of the form $\text{symbol} \rightarrow \text{symbol}_1 \text{symbol}_2 \dots \text{symbol}_n$, where the left argument can be replaced by the sequence of symbols contained in the right argument. Some productions are *terminal*, meaning that they cannot be replaced any longer, while the others are *non-terminal*. Terminal symbols can only appear on the right side, while non-terminals can appear on both sides. Formally a context free grammar is defined as follows

Definition 2.7. A *context-free grammar* is made of the following elements:

- A set of non-terminal symbols N .
- A finite set of terminal symbols Σ , called *alphabet*.
- A non-terminal symbol $S \in N$ called *starting symbol*.
- A set of productions P in the form $N \rightarrow (N \cup \Sigma)^*$.

Note that Definition 2.7 allows *context-free grammars* to process also regular expression, thus context-free grammars are more expressive then regular expressions. In what follows we assume that the terminal symbols are treated as tokens with regular expressions that can be processed by a lexer, but in general a context-free grammar does not require a lexer DFA to process terminal symbols.

In order to check if a sentence is valid in the grammar defined for a language, we perform a process called *derivation*: starting from the symbol

S of the grammar, we recursively replace non-terminal symbols with the right side of their production. The derivation can be done in different ways: we can start expanding the leftmost non-terminal in the production or the rightmost one. The result of the derivation usually generates a data structure called *parse tree* or *abstract syntax tree*, which connects a non-terminal symbol to the symbols obtained through the derivation; the leaves of the tree are terminal symbols.

2.3.1 LR(k) parsers

Simple grammars can be parsed by using *left-to-right parse*, *leftmost-derivation*, *k-tokens lookahead*, meaning that the parser processes a symbol by performing a derivation starting from the leftmost symbol of the production, and looking at the first k tokens of a string of the language. The weakness of this technique is that the parser must predict which production to use only knowing the first k tokens of the right side of the production. For instance, consider the two expression

$$(15 * 3 + 4) - 6$$

$$(15 * 3 + 4)$$

and the grammar

$$\begin{aligned} S &\rightarrow E \text{ eof} \\ E &\rightarrow E + T \\ E &\rightarrow E - T \\ E &\rightarrow T * F \\ E &\rightarrow T / F \\ E &\rightarrow T \\ T &\rightarrow F \\ F &\rightarrow id \\ F &\rightarrow num \\ F &\rightarrow (E) \end{aligned}$$

In the first case the parser should use the production $E \rightarrow E - T$ while in the second it should use the production $E \rightarrow T$. This grammar cannot be parsed by a LL(k) parser because it is not possible to decide which of the two productions must be used just by looking at the first k leftmost tokens. Indeed expressions of that form could have arbitrary length and the lookahead is, in general, insufficient. In general LL(k) grammars are context-free, but not all context-free grammars are LL(k), so such a parser is unable to parse all context-free grammars.

A more powerful parser is the *left-to-right parse*, *rightmost-derivation*, *k-tokens lookahead* or LR(k). This parser maintains a *stack* and an *input* (which is the sentence to parse). The first k tokens of the input are the *lookahead*. The parser uses the stack and the lookahead to perform two different actions:

- *Shift*: The parser moves the first input token to the top of the stack.
- *Reduce*: The parser chooses a grammar production $N_i \rightarrow s_1 s_2 \dots s_j$ and pop s_j, s_{j-1}, \dots, s_1 from the top of the stack. It then pushes N_i at the top of the stack.

The parser uses a DFA to know when to apply a shift action or a reduce action. The DFA is insufficient to process the input, as DFA's are not capable of processing context-free grammars, but it is applied to the stack. The DFA contains edges labelled by the symbols that can appear in the stack, while states contain one of the following actions:

- s_n : shift the symbol and go to state n .
- g_n : go to state n .
- r_k : reduce using the production k in the grammar.
- a : accept, i.e. shift the end-of-file symbol.
- *error*: invalid state, meaning that the sentence is invalid in the grammar.

The automaton is usually represented with a tabular structure, which is called *parsing table*. The element $p_{i,s}$ in the table represents the transition from state i when the symbol at the top of the stack is s .

In order to generate the parsing table (or equivalently the DFA for the parser) we need two support functions, one to generate the possible states the automaton can reach by using grammar productions, and one to generate the actions to advance past the current state. We introduce an additional notation to represent the situation where the parser has reached a certain position while deriving a production.

Definition 2.8. An *item* is any production in the form $N \rightarrow \alpha.X\beta$, meaning that the parser is at the position indicated by the dot where X is a grammar symbol.

At this point we are able to define the *Closure* function, that adds more items to a set of items when the dot is before a non-terminal symbol, which is shown in Algorithm 2.3. Note that, for brevity, we present the version to generate a LR(0) parser, for a LR(1) parser a minor adjustment must be made.

Algorithm 2.3 Closure function for a LR(0) parser

```

function CLOSURE( $I$ )
  repeat
    for all  $N \rightarrow \alpha.X\beta \in I$  do
      for all  $X \rightarrow \gamma$  do
         $I \leftarrow I \cup \{X \rightarrow .\gamma\}$ 
      end for
    end for
  until  $I' \neq I$ 
  return  $I$ 
end function

```

The algorithm starts with an initial set of items I and adds all grammar productions that contain X as left argument as items with the dot at the beginning of their right argument, meaning that the symbols of the production must still be completely parsed.

Now we need a function that, given a set of items, is able to advance the state of the parser past the symbol X . This is shown in Algorithm 2.4.

Algorithm 2.4 Goto function for a LR(0) parser

```

function GOTO( $I, X$ )
   $J \leftarrow \emptyset$ 
  for all  $N \rightarrow \alpha.X\beta \in I$  do
     $J \leftarrow J \cup \{N \rightarrow \alpha.X.\beta\}$ 
  end for
  return CLOSURE( $J$ )
end function

```

The algorithm starts with a set of items and a symbol X and creates a new set of items where the parser position has been moved past the symbol X . It then compute the closure of this new set of items and returns it.

We can now proceed to define the algorithm to generate the LR(0) parser, which is shown in Algorithm 2.5. The initial state is made of all the productions where the left side is the starting symbol, which is equivalent to compute the closure of $S' \rightarrow .S \text{ eof}$. It then proceeds to expand the set of states and the set of actions to perform. Note that we never compute $\text{GOTO}(I, \text{eof})$ but we simply generate an *accept* action. Now, for all actions in E where X is a terminal, we generate a shift action at position (I, X) , for all actions where X is non-terminal we put a goto action at position (I, X) , and finally for a state containing an item $N_k \rightarrow \gamma$. (the parser is at the end of the production) we generate a r_k action at (I, Y) for every token Y .

In general parsing tables can be very large, for this reason it is usually wise to implement a variant of LR(k) parsers called LALR(k) parsers, where all states that contain the same actions but different lookaheads are merged into one, thus reducing the size of the parsing table. LR(1) and LALR(1) parsers are very common, since most of the programming languages can be defined by a LR(1) grammar. For instance, the popular family of parser generators Yacc produces LALR(1) parsers.

Algorithm 2.5 LR(0) parser generation

```

 $T \leftarrow \text{CLOSURE}(\{S' \rightarrow .S \text{ eof}\})$ 
 $E \leftarrow \emptyset$ 
repeat
   $T' \leftarrow T$ 
   $E' \leftarrow E$ 
  for all  $I \in T$  do
    for all  $N \rightarrow \alpha.X\beta \in I$  do
       $J \leftarrow \text{GOTO}(I, X)$ 
       $T \leftarrow T \cup \{J\}$ 
       $E \leftarrow E \cup \{I \xrightarrow{X} J\}$ 
    end for
  end for
until  $E' = E$  and  $T' = T$ 

```

2.3.2 Parser generators

The process of creating a parser can be automated by using a *Parser Generator*. A Parser generator is a programming language that accepts the definition of the grammar of a language and generates a parser (and a lexer) for it. As programming languages generally have a LALR(1) grammar [8], most of parser generators produce a LALR(1) parser. Since in this research work we used F# as a development language, in this section we present the F# lexer and parser generators, belonging to the Yacc generator family, known as *FsLex* and *FsYacc*.

Definition of a lexer in FsLex

FsLex allows to define the tokens with the regular expression syntax. Each FsLex program begins with a header, where the programmer can specify auxiliary modules and functions to use in the lexer. After the header, it is possible to specify relevant regular expressions that are used by the lexer to analyse the tokens with the standard let-binding syntax of F#. The right argument of this binding is a regular expression, which can be composed with the combinators for regular expressions seen in Section 2.2.1. For example the following regular expression can define the syntax for variable names in a programming language:

```
let simpleId = ['a'-'z' 'A'-'Z'] ['a'-'z' 'A'-'Z' '_' '0'-'9']+
```

Regular expression bindings can be used as alias in the lexer definition. A lexer definition is identified by the keyword **rule** for the binding. The right side of a lexer definition contains a call to the function **parse**, which tries to execute one of the rules specified below to parse a token. Each lexer rule generates a result which is a token data structure. Token data structures are specified at parser level (see below). For instance, the following is a lexer able to recognize comparison operators:

```
rule comparisonOperators = parse
| "=" { Parser.EQUAL }
| ">" { Parser.GT }
| ">=" { Parser.GEQ }
| "<" { Parser.LT }
| "<=" { Parser.LEQ }
| "<>" { Parser.NEQ }
```

Note that, in order to provide useful information about the lexing phase, we might need to access, for instance, to the position of the lexer (for error reporting) or to get the string read by the lexer (for example to generate literals when reading numbers or strings). This information is provided by the *lexer buffer*, which is a data structure generated automatically by the parser generator. For example, if the token needs to store its row and column position in the file for error reporting, the definition above can be changed in this way (note the use of the header to define the function **range**):

```
{
  module Lexer

    let range (lexbuf : LexBuffer<_>) = lexbuf.EndPos.Line + 1, lexbuf.
      EndPos.Column
}
```

```

rule comparisonOperators = parse
| ">" { Parser.GT (range lexbuf) }
| ">=" { Parser.GEQ (range lexbuf) }
| "<" { Parser.LT (range lexbuf) }
| "<=" { Parser.LEQ (range lexbuf) }
| "<>" { Parser.NEQ (range lexbuf) }

```

Another useful feature of FsLex is the capability of defining recursive lexers. Let us consider the case of skipping multi-line comments: usually such comments are delimited by a start and end symbol, and the comments spread across multiple lines. For example in C++/Java/C# a multi-line comment is delimited by the symbols `/* */`. The lexer must detect the left delimiter of the multi-line comment, and then keep skipping all the symbols until it detects the right delimiter. This means that the lexer must call itself multiple times, using different lexing rules: one to detect the left delimiter, one to handle new lines or characters inside the comment, and one to detect the right delimiter. Furthermore, after handling the comment, the lexer must go back to processing the program normally. The following code shows how to implement such lexer:

```

{
  module Lexer

    let newline (lexbuf : LexBuffer<_>) = lexbuf.EndPos <- lexbuf.EndPos.
      NextLine
    let range (lexbuf : LexBuffer<_>) = lexbuf.EndPos.Line + 1, lexbuf.
      EndPos.Column
  }
  let newline = ('\n' | '\r' '\n')

  rule comment = parse
  | "/*" { programTokens lexbuf }
  | newline { newline lexbuf; comment lexbuf }
  | _ { comment lexbuf }

  and programTokens = parse
  | "/*" { comment lexbuf }
  ...
  //other token definitions

```

Note that `programTokens` calls `comment` when it detects the left delimiter of a multi-line comment. `comment` keeps calling itself until the right delimiter is detected, where it jumps back to `programTokens`.

Definition of a parser in FsYacc

FsYacc allows to define the grammar of a language in terms of productions of a context free grammar. As for the lexer, the parser definition starts with a header where the programmer can specify custom code and modules to use. The grammar defines terminal symbols as tokens, identified by the keyword `%token`. A token specified the name to be used in the grammar productions, and a series of type parameters that are used to store data in a token. For example, the following tokens might be used in a parser for arithmetic expressions:

```

%token PLUS MINUS MUL DIV LPAR RPAR
%token <double> NUMBER

```

Whenever a terminal symbol is encountered during the parsing phase, the parser calls the lexer to generate the data structure for the token. The lexer tries to match the string provided by the parser by using one of its rule and, if it succeeds, it returns the appropriate token data structure. In this part of the grammar we must also specify the starting symbol. This symbol is defined through the keyword `%start`. Since usually we want to generate an abstract syntax tree for the grammar (which must be manually defined), we can specify a return type generated by the parser with the keyword `%type`. For an arithmetic expression this would be, for instance

```
%start start
%type <Expr> start
```

In this section it is also possible to define the operators associativity and precedence, through the keywords `%left`, `%right`, and `%nonassoc`. Terms defined in the same associativity line have the same precedence, and the precedence is ordered according to the line number, so if a term associativity is defined below another, it has higher precedence.

After the terminal symbol definitions, the grammar must specify productions. A production is defined in the following way:

```
productionName:
| rule_1 { action_1 }
| rule_2 { action_2 }
...
| rule_n { action_n }
```

Each action defines the code that the parser executes when that rule is matched. Usually this part is used to build the nodes of the syntax tree, but there is no restriction in what the action can perform, as long as it is valid F# code. It is possible to access the result of evaluating a term in the production by using an index preceded by the symbol `%`, where `%1` refers to the first term in the right hand-side of the production. For example this code might be used to parse an arithmetic expression:

```
%{
  open AST
}%

%token PLUS MINUS MUL DIV LPAR RPAR EOF
%token <float> NUMBER

%left PLUS MINUS
%left MUL DIV

%start start
%type <Expr> start

start : Expression EOF { %1 }

Expression:
| NUMBER { Number %1 }
| Expression PLUS Expression { Plus(%1,%3) }
| Expression MINUS Expression { Minus(%1,%3) }
| Expression MUL Expression { Mul(%1,%3) }
| Expression DIV Expression { Div(%1,%3) }
| LPAR Expression RPAR { Nested(%2) }
```



```

module AST

type Expression =
| Number of float
| Plus of Expression * Expression
| Minus of Expression * Expression
| Mul of Expression * Expression
| Div of Expression * Expression
| Nested of Expression

```

2.3.3 Monadic parsers

Monadic parsing is an alternative to traditional parsers, such as LR(k) and LALR(k) presented above. Monadic parsers have inferior performance with respect to LR(k) and LALR(k) [29] parsers but they are extensible, i.e. they do not rely on a limited set of combinators to describe the grammar of a language as for parser generators. Monadic parsers were extensively explained in [29, 43], here we present a variation that can deal also with error handling. Before explaining how to implement a monadic parser, we introduce the concept of Monad:

Definition 2.9. A *Monad* is a tern made of the following elements:

- A type constructor M .
- A unary operation $Return :: a \rightarrow M a$.
- A binary operation $Bind :: M a \rightarrow (a \rightarrow M b) \rightarrow M b$. The bind can also be written by using the symbol $>>=$.

where both operations satisfy the following properties:

- $a >>= return \equiv a$.
- $(a >>= f) >>= g \equiv a >>= (\lambda x. fx >>= g)$.

We now proceed to define a parser monad by defining (i) the type constructor for the parser, (ii) the unary operator, (iii) the binary operator, and (iv) parser combinators as an example of the extensibility of the parser monad. Note that below we provide an implementation in F#, which does not have type classes as Haskell, so the parser monad does not use any type argument and directly defines the operators for this specific instance of monad.

Parser type constructor and monadic operations

A parser is defined in literature as a function that takes as input a text and returns a list of pairs made of the parsing result and the rest of the text to process. The parsing result is usually the syntax tree generated by the parser. The result is a list because the same syntactical structure might be processed in different ways. By convention, an empty list denotes a parser failure. Here we propose a variation of this traditional implementation in order to provide a better error report.

In this alternative implementation, the parser is a function that takes as input the text to process, a *parsing context* that might hold auxiliary information necessary for the parsing, the current position of the parser in the text, and returns either a tuple containing the parsing result, the text

left to process, an updated context, and the updated position, or an error in case of a parser failure.

```
type Parser<'a, 'ctxt> = { Parse : List<char> -> 'ctxt -> Position ->
    Either<'a * List<char> * 'ctxt * Position, Error>}

static member Make(p:List<char> -> 'ctxt -> Position -> Either<'a * List
    <char> * 'ctxt * Position, Error>) : Parser<'a,'ctxt> = { Parse = p
    }
```

The *return* operation should take as input a generic value of type 'a and return a `Parser<'a,'ctxt>`. The return simply creates the parser function for the given input:

```
member this.Return(x:'a) : Parser<'a,'ctxt> =
    (fun buf ctxt pos -> First(x, buf, ctxt, pos)) |> Parser.Make
```

According to the Definition 2.9, the bind operator must take as input a `Parser<'a>`, a function `'a -> Parser<'b>` and return `Parser<'b>`. The bind generates a function that runs the input parser on the text. The result of the input parser can, according to its definition, contain a parsing result or an error in case of failure. The function generated by the bind must be able to handle these two situations: in case of a correct result the function creates a new parser using the parsing result and runs it on the remaining portion of the text, while in case of an error it simply outputs the error. In this way, when parsing fails, the error will be propagated ahead.

```
member this.Bind(p:Parser<'a,'ctxt>, k:'a->Parser<'b,'ctxt>) : Parser<'b
    ,'ctxt> =
    (fun buf ctxt pos ->
        let all_res = p.Parse buf ctxt pos
        match all_res with
        | First pires ->
            let res, restBuf, ctxt', pos' = pires
            (k res).Parse restBuf ctxt' pos'
        | Second err -> Second err ) |> Parser.Make
```

Parser combinators

With the parser monad implemented above, we can implement several parser combinators that can be used to define the grammar of a language. Here we show only a small glimpse of the possible combinators that can be implemented.

The first parser combinator that we present is the *choice*. The choice takes as input two parsers and runs the first. If the first parser succeeds than its result is returned, otherwise the second is run. If it succeeds its result is return, otherwise the whole parser outputs an error. This combinator is useful, for instance, when there might be two possible choices for a token in a statement. For instance, in either Java or C# is possible to exchange the order of the access modifier and the static modifier in the method declaration, thus both `public static` or `static public` are valid combinations. This combinator would try to parse the declaration in the first way, and if it fails it will try also the second option. Of course if the syntax of both combinations is wrong the parser will fail completely. The code for the combinator is shown below:

```
static member (++) (p1:Parser<'a','ctxt>, p2:Parser<'a','ctxt>) : Parser<'
a','ctxt> =
    (fun buf ctxt p ->
        match p1.Parse buf ctxt p with
        | Second err1 ->
            match p2.Parse buf ctxt p with
            | Second err2 -> Second err2
            | p2res -> p2res
        | p1res -> p1res) |> Parser.Make
```

A useful variation of this combinator, is the one that executes two parsers with different generic types and returns a **Either** data type, containing either the result of the first or the second.

```
static member (+) (p1:Parser<'a','ctxt>, p2:Parser<'b','ctxt>) : Parser<
Either<'a','b','ctxt> =
    (fun buf ctxt p ->
        match p1.Parse buf ctxt p with
        | Second err1 ->
            match p2.Parse buf ctxt p with
            | Second err2 -> Second(err2)
            | First p2res ->
                let res,restBuf,ctxt',pos = p2res
                First(Second res, restBuf, ctxt', pos)
        | First p1res ->
            let res, restBuf, ctxt', pos = p1res
            First((First res), restBuf, ctxt', pos)) |> Parser.Make
```

Other combinators are possible, but for brevity we have only shown two. It should appear clear how this approach is completely extensible with no limitations. Any combinator would take as input two parsers and define the type of the resulting parser. The implementation will contain the logic to combine two parsers together. For example, another parser combinator is the application of 0 or more times of the same parser.

To complete this discussion, we now show how to parse a specific character and a keyword. The parser for a character takes as input the text to process and the character to match. If the input text is empty of course the parser immediately fails because no character will ever be matched. Otherwise if the first character of the text matches the one provided then we return the matched character as result and the rest of the text to process, otherwise we output an error. The function also takes care of updating the position of the parser accordingly and to skip line breaks.

```
let character(c:char) : Parser<char, 'ctxt> =
    (fun buf ctxt (pos:Position) ->
        match buf : List<char> with
        | x::cs when x = c ->
            let pos' =
                if x = '\n' then
                    pos.NextLine
                else
                    pos.NextCol
            First( c, cs, ctxt, pos')
        | _ ->
            Second (Error(pos, sprintf "Expected character %A" c))) |> Parser.
            Make
```

The word parser takes as input the text to process and the word to match. It then applies the character parser to the word until it has all been processed. In the code below the syntax **let! x = y** is a syntactical sugar for **y >>= fun x -> ...** in the fashion of Haskell **do** notation.

```

let rec word (w:List<char>) : Parser<List<char>, 'ctxt> =
  p{
    match w with
    | x::xs ->
      let! c = character x
      let! cs = word xs
      return c::cs
    | [] ->
      return []
  }

```

2.4 Type systems and type checking

Being able to verify the correctness of a program is a crucial aspect of programming. When dealing with low-level languages is generally difficult to verify and grant the correctness of a program since a language such assembly does not provide abstractions for the purpose. Modern high-level programming languages, on the other hand, generally provide a way to type their constructs. A type system is a syntactic method that assigns a property called *type* to the constructs of a programming language, in order to prove that a program does not have certain unwanted behaviours [36]. Type systems are generally expressed in the form of inference rules [14, 36], made of a set of premises, that must be verified in order to assign to the language construct the type defined in the conclusion. An inference rule is a logical rule in the form:

$$\begin{array}{c}
 \text{premise}_1 \\
 \text{premise}_2 \\
 \vdots \\
 \text{premise}_n \\
 \hline
 \text{conclusion}
 \end{array}$$

where all the premises must be true in order to evaluate the conclusion. Usually the type rules make use of a *typing environment*, which is an association between language constructs and types. For example the following rule defines the typing of an **if-then-else** and a **while-do** statement in an imperative language.

$$\frac{\Gamma \vdash c : \text{bool} \quad \Gamma \vdash t \quad \Gamma \vdash e}{\Gamma \vdash \text{if } c \text{ then } t \text{ else } e}$$

$$\frac{\Gamma \vdash c : \text{bool} \quad \Gamma \vdash w}{\Gamma \vdash \text{while } c \text{ do } w}$$

In these rules Γ is the environment. The type rule first evaluates the premises, which means that if the condition of the **if-then-else** has type `bool` and the evaluation of the **then** and **else** block succeeds, then the whole

if-then-else is correctly typed. Analogously, for the **while-do**, if the condition has type `bool` and the evaluation of the while block is correctly typed, then the whole **while-do** is correctly typed. Note that control structures code blocks are usually not given a type, rather they are considered correct if all their statements are correctly typed. An equivalent way of expressing this is using a special type called *unit* for constructs that do not return a value. This expedient is widely used in hybrid functional programming languages such as F# or CamL. The equivalent rules for the construct above would be:

$$\frac{\Gamma \vdash c : \text{bool} \quad \Gamma \vdash t : \text{unit} \quad \Gamma \vdash e : \text{unit}}{\Gamma \vdash \text{if } c \text{ then } t \text{ else } e : \text{unit}}$$

$$\frac{\Gamma \vdash c : \text{bool} \quad \Gamma \vdash w : \text{unit}}{\Gamma \vdash \text{while } c \text{ do } w : \text{unit}}$$

Typing a construct of the language requires to evaluate its corresponding typing rule. Unlike for parsers, there exist no tools capable of automatically generate a type checker given the type rules definition, thus the behaviour of each type rule must be implemented in the host language in which the compiler is defined. Independently of the chosen language, the behaviour will always be the following : (i) evaluate a premise, (ii) if the evaluation of the premise fails then the construct fails the type check and an error is returned, (iii) repeat step 1 and 2 until all the premises have been evaluated, and (iv) assign the type to the construct that is defined in the rule conclusion.

During this process the compiler generates a data structure called *symbol table*, which contains information about the type checking process and maintains the type environment.

At the end of the type check, the program is correct with respect to types. At this point, depending on the chosen target language, the compiler might discard or keep the information about the typing process. Usually when targeting a high-level programming language, the information about the types is kept because they are necessary during the code generation process in order to, for instance, generate the proper variable declaration statements. On the other hand, when targeting a low-level untyped programming language, such as assembly, the type information can be discarded.

2.5 Semantics and code generation

Semantics define how the language abstractions behave and can be expressed in different ways, for example with a term-rewriting system [31], reduction semantics [22] or with the operational semantics [24]. Below we provide a description of these three possible representations for the semantics:

Term-rewriting semantics Term-rewriting semantics define a set of rewriting rules that take as input a construct of the language and define how to rewrite it into another form. The rewriting process usually ends when a rewrite rule is replaced by itself. For example the **if-then-else** statements and **while-do** statements can be rewritten with these rules (the ; symbol denotes a sequence of statements):

```

if true then t else e ; k → t;k
if false then t else e ; k → e;k
while (c = true) do w ; k → w ; while (c) do w ; k
while (c = false) do w ; k → k

```

Reduction semantics Reduction semantics use a *reduction context*, which is a program or a fragment of program with a *hole* (denoted by the symbol \square) as placeholder to mark where the next computational step is taking place. For example the **if-then-else** and **while-do** statements semantics can be represented in the following way¹:

```

if  $\square$  then s1 else s2 ; k
if true then t else e → t
if false then t else e → e
while  $\square$  do w
while true do w → w ; while  $\square$  do w
while false do w → skip

```

Operational semantics Operational semantics define the behaviour of language constructs in terms of logical rules similar to those used for type systems. For instance, the **if-then-else** and **while-do** semantics are expressed as

$$\frac{\langle c \rangle \Rightarrow \text{true}}{\langle \text{if } c \text{ then } T \text{ else } E ; k \rangle \Rightarrow T ; k}$$

$$\frac{\langle c \rangle \Rightarrow \text{false}}{\langle \text{if } c \text{ then } T \text{ else } E ; k \rangle \Rightarrow E ; k}$$

$$\frac{\langle c \rangle \Rightarrow \text{true}}{\langle \text{while } c \text{ do } L ; k \rangle \Rightarrow L ; \text{while } c \text{ do } L ; k}$$

$$\frac{\langle c \rangle \Rightarrow \text{false}}{\langle \text{while } c \text{ do } L ; k \rangle \Rightarrow k}$$

Regardless of the formal representation chosen for the semantics, this must be encoded in the abstractions of the target language during the code generation phase. When choosing a high-level target language encoding the operational semantics of a similar high-level language might be trivial,

¹`skip` is a statement that simply skips to the next statement in a sequence of statements

but generating for instance the code for a functional programming language into an imperative language might prove difficult. For this reason, the code generation step might be preceded by an intermediate code generation step. The intermediate language is usually a simple programming language close to the target language. Notable examples of this are the *three-address code* [7], and the intermediate language used in the *Glasgow Haskell Compiler* [28] and *Utrecht Haskell Compiler* [20].

2.6 Metaprogramming and metacompilers

This section aims to provide the reader with sufficient information to understand the concept of metaprogramming. In this section we explain what metaprogramming is and present existing metacompilation approaches existing in scientific literature. We start by defining what metaprogramming is and we present techniques of metaprogramming in existing programming languages. We then proceed by presenting how different existing metacompilers work.

Metaprogramming is the process of writing computer programs with the ability to treat programs as their data [17]. Metaprogramming takes as input a program written in a meta-language to define a programming language called *object language*, a program written in the object language, and outputs executable code able to run the program. Metaprogramming can be achieved in two different ways: (i) by using opportune language abstractions provided by a general-purpose programming language, or (ii) using a dedicated metacompiler. In what follows we provide examples in both areas.

2.6.1 Template metaprogramming

Template metaprogramming uses class templates to operate on numbers and types as data. In this section we provide examples in C++ templates, but other languages allow template metaprogramming, with notable examples being Lisp macros and Haskell templates [39]. The template language uses template recursion as loop construct and template specialization as decisional construct.

To better understand how this works, we will implement the factorial function with templates. It is well known that, by definition, the factorial of 0 is 1, while the factorial of a number n is n multiplied by the factorial of $n - 1$. Our template meta-program will thus contain two templates: one for the base case of the recursion and one for the recursive step. The base case of the recursion uses template specialization to stop the computation and immediately return 1:

```
template<>
struct Factorial<0>
{
    enum { RET = 1 };
};
```

This template contains an enumeration type whose only value is 1. The recursive step will take as input a generic template parameter and recursively call the template definition.

```
template<int n>
struct Factorial
{
    enum { RET = Factorial<n - 1>::RET * n };
};
```

When the `Factorial` template is instantiated with a value different from 0, the non-specialized version is used by the C++ compiler. The enumeration case `RET` then gets the value of `RET` for the same template instantiated for `n - 1` and multiplied by `n`. The generation of templates and their enumeration cases will stop when the template instantiation will be invoked with `Factorial<0>`, which will use the specialized version. Note the use of the scope resolution operator to access the value of the enumeration case. This template can be used instantiating the template with an integer constant, for instance:

```
int main()
{
    cout << Factorial<5>::RET << endl;
}
```

Note that template instantiation is performed at compile-time, so the result of the factorial is actually inlined by the compiler every time the template is instantiated.

A more interesting example is about how to define recursive data structures. Let us consider the implementation of lists in a functional programming language:

```
type List<'a> =
| Empty
| Cons of 'a * List<'a>
```

where the list `[3,4,5,6]` can be built as `Cons(3,Cons(4,Cons(5,Cons(6,Empty))))`. This list representation can be defined with template metaprogramming by defining a template specialization for the empty list, and a non-specialized template for a non-empty list.

```
struct NIL
{
    typedef NIL Head;
    typedef NIL Tail;
};

template<class T, class Tail_ = NIL>
struct Cons
{
    typedef T Head;
    typedef Tail_ Tail;
};
```

Note that the assignment in the template definition specifies an optional template parameter, in the same way as optional method arguments.

Now let us try to define a function to calculate the length of an arbitrary list. This function will have as a base case the empty list, for which it returns 0, otherwise it returns 1 plus the length of the tail. Again this can be implemented with a specialized template and a non-specialized template.


```

template<class List>
struct Length
{
    static const unsigned int RET = Length<List::Tail>::RET + 1;
};

template<>
struct Length<NIL>
{
    static const unsigned int RET = 0;
};

```

The first template recursively class the `Length` template with the type of the tail of the list, extracts the `RET` field and adds 1. The second template is a specialization created with the `NIL` type and immediately sets the field `RET` to 0.

In order to test this function we must create a template for a data type (which is a meta-data) that we want to store in the list. In this example we show how to test the function for a list of integers. First of all we must create a template for an integer:

```

template<int n>
struct Int
{
    static const int Value = n;
};

```

This is necessary in order to be able to store the values of the list elements. At this point, the list can be created by calling `Cons` and passing the `Int` data type as argument and. `Length` can then be used with the type of the list that has been created and then we can access its `RET` field.

```

typedef Cons<Int<3>, Cons<Int<4>, Cons<Int<5>, Cons<Int<6>>>> testList;
cout << Length<testList>::RET << endl;

```

Template metaprogramming complexity can grow exponentially, for example when we want to get the values of the list elements, to the point that a simple function as `nth` requires several templates. We omit the details here, but the reader can find additional information in Appendix A.

2.6.2 Metacompilers

Metacompilers are a special class of compilers used to implement other compilers. A metacompiler takes as input the definition of the syntax, semantics, and possibly the type system of the object language, a program written in the object language, and outputs executable code for it. Metacompilers are written either in a general purpose programming language or in their own meta-language through the process of self-hosting. Self-hosting compilation requires to write a prototypical version of the compiler in another language or an interpreter for it and then use it to compile the implementation of a subsequent version. In this section we present some examples of existing meta-compilers.

META-II

META-II is one of the earliest metacompilers and, for this reason, quite limited in its capabilities. META-II allows to express the syntax of the object

language and actions for the code generation. A meta-program in META-II is made of grammatical symbols, meta-variables, and equations that define the terms of the grammar. A symbol is written as a string surrounded by quotes and beginning with a period, a meta-variable is a string starting with a alphabetical character and followed by an arbitrary amount of alphanumeric characters, and an equation is a sequence of consecutive symbols or ids to indicate concatenation. Alternation is defined with the symbol /, which can be used together with the keyword .EMPTY to define alternation. For instance

```
BOOLEAN = '.TRUE' / '.FALSE'
```

defines boolean literals. The meta-language is able to recognize built-in symbols such as identifiers, denoted with .ID, strings represented by .STRING, and numbers represented by .NUMBER. These are to be intended as identifiers, strings, and numbers in the object language. For example the family of expressions:

```
A
A + B
A + B * C
(A + B) * C
```

can be encoded by the following equations in META-II

```
EX3 = .ID / '(' EX1 ')',
EX2 = EX3 ('*' EX2 / .EMPTY),
EX1 = EX2 ('+' EX1 / .EMPTY)
```

Sequences (as in the Kleene closure for regular expressions) can be expressed using the symbol \$. For example

```
SEQA = $ 'A',
```

represents a sequence containing the letter A. META-II allows to associate actions to equation for the code generation. Each action generates assembly code for an interpreter called META-II machine, which is able to execute it. The action of code generation is marked with the keyword .OUT, for instance

```
EX3 = .ID .OUT('LD ' *) / '(' EX1 ')'
```

generates the literal output and the special symbol found in EX3.

META-II is a self-hosting compiler, i.e. it is implemented in META-II itself.

RML

RML [35] (*Relational Meta-Language*) uses a meta-language based on operational semantics. A program in RML consists of data definitions in a syntax similar to CamL discriminate unions, and relations containing axioms and inference rules. An axiom is an inference rule without premises, while an inference rule generates an output if the premises correctly evaluate. For example the following snippet defines the data type for an arithmetic expression and a symbol table for the evaluation.

```

datatype Expr =
| INT of int
| VAR of string
| ADD of Expr * Expr

type Env = (string * int) list

```

Axioms and inference rules are grouped together into a *relation*. For example, the following relation can be used to evaluate an arithmetic expression:

```

relation eval =
  axiom eval(env, INT i) => i

  rule
    lookup(env, x) => i
    -----
    eval(env, VAR x) => i

  rule
    eval (env, left) => i1 &
    eval (env, right) => i2
    i1 + i2 => v
    -----
    eval (env, ADD(left,right)) => v

```

During the code generation phase, each relation is translated into a first-order logic representation, which consists of a series of `match` structures that check the structure of the arguments. For example the rule above would be translated into:²

```

(and (match [(arg1 env)
              (arg2 ADD(left,right))]))
(and (call eval [env left] [result1]))
(and match [result1 i1])
(and (call eval [env right] [result2]))
(and match [result2 i2])
(and call [i1 + i2] [result3])
(and match [result3 v])
(return v)

```

This code is later translated into a continuation-passing style form, which is later generated as C code. The compiler performs heavy optimization on tail calls generated code through the use of a technique called *dispatching switches*.

Stratego

Stratego [13] is a metacompiler that uses a term-rewriting semantics as meta-language to define its programs. A stratego program consists of a series of terms in the form

$$t := c(t_1, t_2, \dots, t_n)$$

where c is a constructor that accepts n other terms as arguments. The syntax of Stratego has been enriched with additional syntax to handle “traditional” data structures, such as string, integer, float, constants, and lists:

²We use a prefix notation in Lisp style

$$pt := s \mid i \mid f \mid [t_1, t_2, \dots, t_n] \mid (t_1, t_2, \dots, t_n) \mid c(t_1, t_2, \dots, t_n)$$

Terms can be extended with a list of annotations that are terms themselves:

$$t := pt \mid pt \{t_1, t_2, \dots, t_n\}$$

Stratego requires that the meta-program specifies the signature of term constructors. For example simple arithmetic expressions can be defined as

```
signature
  sorts Id Expr
  constructors
    Var : Id -> Expr
    Plus : Expr * Expr -> Expr
```

Note that Stratego is an untyped language, so types are not statically checked and the compiler only checks that constructors are declared and have the correct arity.

Rewrite rules define how terms are evaluated, for example the following is a rewrite rule to evaluate a binary operator in an arithmetic expression:

```
EvalBinOp : Plus(Int(i), Int(j)) -> Int(k) where k := <add>(i, j)
```

Note that rewrite rules support conditionals, i.e. in the rule above we are able to specify that `k` is the result of adding the numbers `i` and `j` given as arguments.

Stratego compiler is a self-hosting compiler, meaning that the Stratego meta-language is defined in the meta-language itself. A first version of Stratego was written in SML, which was then re-used to compile a further iteration written in Stratego. Stratego compiles programs to C, where the code generation transformations were expressed in Stratego itself.

Chapter 3

Metacasanova

Typing is no substitute for
thinking

Dartmouth Basic manual, 1964

This chapter aims to provide the reader with additional motivation about the advantage of meta-compilers and sufficient details on the Metacasanova compiler architecture. In Section 3.1 we show that there exists a common pattern when implementing the formal definition of the type system and semantics of a programming language into the abstraction of a general-purpose programming language. In Section 3.2 we define the requirements of Metacasanova and give an informal overview about the structure of a meta-program and we provide a formalization of its semantics. In Section 3.3 we explain the working principles of the meta-compiler and the involved compilation stages. In Section 3.4 we present the details of the Metacasanova grammar and parser; we then explain the subsequent parsing post-processing phase and how the post-processor re-processes the generated AST. In Section 3.5 we explain how the type checking of a meta-program is performed and in what cases it fails. In Section 3.6 we explain how the abstractions of Metacasanova are mapped into the abstractions of the C# target code.

3.1 Repetitive steps in compiler development

In Chapter 2 we gave an overview of the necessary steps involved in developing a compiler. We showed that the lexing/parsing phase is simple enough to be automated using a lexer/parser generator. Such software takes as input the grammar and the definitions of regular expressions to define the tokens of the language and produces output code containing that is able to parse a program written in a programming language defined by that grammar. However, the steps involved in the following phases, namely the *type checking* and *operational semantics* implementation follow a recurring pattern, but in general the behaviour of the type system and the code generation reflecting the behaviour of the operational semantics must be hard-coded in

the host language in which the compiler is being implemented. Below we present two examples to show how these behaviours can be implemented in two different general purpose programming languages and show that both follow the same pattern.

3.1.1 Hard-coded implementation of type rules

As shown in Section 2.4, type rules are expressed in the form of logical rules. Let us consider the type rules for the **if-then-else** and **while-do** statements presented in Section 2.4 in the version that assigns the type *unit* to the code blocks for convenience. In a programming language that supports discriminated unions as a language abstraction (like Haskell or F#), the syntactical element in the abstract syntax tree of the language can be expressed as

```
type Statement =
| If of Expr * List<Statement> * List<Statement>
| While of Expr * List<Statement>
... //other statements
```

The type checking of the **if** statement requires that to check that the condition has type **bool** and that both code blocks have type **unit** (or **void**). The type checking of the **while-do** is analogous, except only one code block is used. We can then define a function **eval** that, given the environment (here we call it *symbol table*) and a statement as input, returns the type given by the rule or an error if all type rules for that statement fail to correctly evaluate. For the **if-then-else** the implementation is the following:

```
let rec evalStmt (symbolTable : SymbolTable) (stmt : Statement) : Type =
match stmt with
... //other statements
| If (condition, _then, _else) ->
    let conditionType = evalExpr symbolTable condition
    let thenType = evalStmt symbolTable _then
    let elseType = evalStmt symbolTable _else
    if conditionType <> Boolean then
        failwith "Invalid condition type"
    elif thenType <> Unit then
        failwith "The type of then must be unit"
    elif elseType <> Unit then
        failwith "The type of else must be unit"
    else
        Unit
... //other statements
```

The function first executes pattern matching on the statement to identify the correct inference rule to use during the typing. It then proceeds to evaluate the premises (type of the condition and of the statement blocks) and to check their result. If all premises evaluate successfully the type contained in the conclusion is returned. Note that the function **evalExpr** is a function able to evaluate the type rule for expressions and return their type. The implementation of the **while-do** follows the same logic:

```
let rec evalStmt (symbolTable : SymbolTable) (stmt : Statement) : Type =
match stmt with
... //other statements
| While (condition, _do) ->
    let conditionType = evalExpr symbolTable condition
```

```

let doType = evalStmt symbolTable stmt
if conditionType <> Boolean then
  failwith "Invalid condition type"
elif doType <> Unit then
  failwith "The type of the do block must be unit"
else
  Unit
... //other statements

```

In languages that do not provide abstractions such as discriminated unions, the type for statements must exploit polymorphism to implement the same behaviour. A statement will be represented as an interface exhibiting the behaviour of a visitor pattern:

```

public interface Statement
{
  Type Visit(StatementVisitor visitor);
}

public interface StatementVisitor
{
  ... //other statements
  Type OnIf(Expression condition, List<Statement> _then, List<Statement>
    _else);
  Type OnWhile(Expression condition, List<Statement> _do);
  ... //other statements
}

```

The behaviour of the inference rule for the **if-then-else** statement is modelled by a class implementing the **StatementVisitor** interface. This class contains a method **OnIf** that implements the behaviour of the type rule itself.

```

public class StatementEvaluator : StatementVisitor
{
  ... //evaluation of other statements
  public Type OnIf(Expression condition, List<Statement> _then, List<
    Statement> _else)
  {
    Type conditionType = condition.visit(new ExpressionEvaluator());
    Type thenType = _then.Visit(new StatementEvaluator());
    Type elseType = _else.Visit(new StatementEvaluator());
    if (!conditionType.Equals(new Boolean()))
    {
      throw new TypeException("Invalid condition type");
    }
    else if (!thenType.Equals(new Unit()))
    {
      throw new TypeException("The type of then must be unit");
    }
    else if (!elseType.Equals(new Unit()))
    {
      throw new TypeException("The type of else must be unit");
    }
    else
    {
      return new Unit();
    }
  }
}

... //evaluation of other statements
}

public class If : Statement
{

```

```

Expression Condition;
List<Statement> Then;
List<Statement> Else;

public Type Visit(StatementVisitor visitor)
{
    return visitor.OnIf(this.Condition, this.Then, this.Else)
}
}

```

Analogously for the **while-do** we have

```

public class StatementEvaluator : StatementVisitor
{
    ... //evaluation of other statements
    public Type OnWhile(Expression condition, List<Statement> _do)
    {
        Type conditionType = condition.Visit(new ExpressionEvaluator());
        Type doType = _do.Visit(new StatementEvaluator());
        if (!conditionType.Equals(new Boolean()))
        {
            throw new TypeException("Invalid condition type");
        }
        else if (!doType.Equals(new Unit()))
        {
            throw new TypeException("The type of do must be unit");
        }
        else
        {
            return new Unit();
        }
    }
    ... //evaluation of other statements
}

public class While : Statement
{
    Expression Condition;
    List<Statement> Do;

    public Type Visit(StatementVisitor visitor)
    {
        return visitor.OnWhile(this.Condition, this.Do);
    }
}

```

Generalization

In general, for a node of the abstract syntax tree (AST) α (like **Statements**) containing syntactical structures σ_i constructed with a certain amount of arguments of type $\epsilon_{\sigma_{i_1}}, \dots, \epsilon_{\sigma_{i_m}}$ (such as the condition or the statement block in a control structure), the general representation of a hard-coded type rule in a language with discriminated unions is obtained by creating a union type α having a case σ_i with arguments $\epsilon_{\sigma_{i_j}}$ for each syntactical element.

```

type  $\alpha$  =
|  $\sigma_1$  of  $\tau_{\sigma_{1_1}} * \dots * \tau_{\sigma_{1_m}}$ 
...
|  $\sigma_n$  of  $\tau_{\sigma_{n_1}} * \dots * \tau_{\sigma_{n_m}}$ 

```


Evaluating the inference rule through an evaluation function requires first to find out which must be applied by matching the pattern of the syntactical structure from the node of the AST. Later, we need to evaluate each of the premises with the appropriate evaluation function: if the result of each evaluation is what the rule expects (for instance, that the condition has type boolean in the **if-then-else**) then we return the result of the evaluation rule contained in the right part of the conclusion.

Let us consider a conclusion $\sigma_j(\epsilon_{\sigma_{j_1}} \dots \epsilon_{\sigma_{j_m}})$ (where each ϵ is one of the arguments used to construct the case of the discriminate union) with a result of the evaluation ρ_{σ_j} a set of premises π_1, \dots, π_k that are evaluated through an evaluation function φ_{π_i} , $i = 1, \dots, k$ returning a result ρ_{π_i} . Let us assume that ρ'_{π_i} is the expected result for the premise evaluated through φ_{π_i} . As usual, Γ defines the environment (symbol table). The type rule that we are trying to execute will thus have the following structure:

$$\frac{\begin{array}{c} \Gamma \vdash \varphi_{\pi_1} \pi_1 : \rho'_{\pi_1} \\ \vdots \\ \Gamma \vdash \varphi_{\pi_i} \pi_i : \rho'_{\pi_i} \\ \vdots \\ \Gamma \vdash \varphi_{\pi_k} \pi_k : \rho'_{\pi_k} \end{array}}{\Gamma \vdash \sigma_j(\epsilon_{\sigma_{j_1}}, \dots, \epsilon_{\sigma_{j_m}}) : \rho_{\sigma_j}}$$

Given the considerations above, the code necessary for the evaluation will be the following:

```
let rec  $\varphi_{\sigma_j}$   $\Gamma$   $\sigma$  =
  match  $\sigma$  with
  ... //other pattern matching expressions for other rules
  |  $\sigma_j(\epsilon_{\sigma_{j_1}}, \dots, \epsilon_{\sigma_{j_m}})$  ->
    let  $\rho_{\pi_1} = \varphi_{\pi_1} \Gamma \pi_1$ 
    .
    .
    let  $\rho_{\pi_i} = \varphi_{\pi_i} \Gamma \pi_i$ 
    .
    .
    let  $\rho_{\pi_k} = \varphi_{\pi_k} \Gamma \pi_k$ 
    if  $\rho_1 \neq \rho'_1$  then
      failwith "Type error"
    .
    .
    elif  $\rho_i \neq \rho'_i$  then
      failwith "Type error"
    .
    .
    elif  $\rho_m \neq \rho'_m$  then
      failwith "Type error"
    else
       $\rho_{\sigma_j}$ 
  ... //other pattern matching expressions for other rules
```

Each evaluation function is recursive because a premise might need to run the same evaluation function (see the example of the statements above). The function contains a pattern matching that selects the correct inference rule to be used for that syntactical structure. For example, in the case of the statements, it will try to match all the possible syntactical structures for the statements and select the correct one for the input; for instance, if we are running the rule for the **if-then-else** then the pattern matching will select the match case for **if-then-else**. Note that γ will surely be matched by one of the match cases because at this point we have a correctly generated AST after the parsing phase.

Each premise runs the appropriate evaluation function and returns a result. This result is compared with the one expected by the inference rule, and if the comparison fails the function reports a type error. If all comparisons succeed, then the result of the conclusion is returned.

In a language that does not provide discriminated unions and pattern matching the generalization is more complex: the abstract syntax tree element must be represented by an interface containing the signature of a method `Visit` used to perform an operation on a specific (polymorphic) syntactical structure. We also require the interface for the visitor pattern with the signature of the functions to run for each polymorphic instance of `Statement`. In this version we assume that the type of the result of the evaluation function for σ_j returns a type $\tau_{\rho\sigma_j}$ (which in the previous version could be omitted thanks to the type inference typical of functional programming languages):

```
public interface  $\alpha$ 
{
    public  $\tau_{\sigma_j}$  Visit(Visitor visitor)
}

public interface Visitor<T>
{
    T  $\varphi_{\sigma_j}(\tau_{\Gamma} \Gamma, \tau_{\sigma_{j_1}} \epsilon_{\sigma_{j_1}}, \dots, \tau_{\sigma_{j_m}} \epsilon_{\sigma_{j_m}})$ ;
    ... //other statements
}
```

Then we have to implement the visitor for the type rule for statements and a class for a specific statement:

```
public class Evaluator : Visitor< $\tau_{\rho\sigma_j}$ >
{
    ... //evaluation of other statements
    public  $\tau_{\rho\sigma_j} \varphi_{\sigma_j}(\tau_{\Gamma} \Gamma, \tau_{\sigma_{j_1}} \epsilon_{\sigma_{j_1}}, \dots, \tau_{\sigma_{j_m}} \epsilon_{\sigma_{j_m}})$ 
    {
         $\tau_{\rho\phi_1} \rho_{\pi_1} = \varphi_{\pi_1}(\Gamma, \pi_1)$ ;
        .
        .
        .
         $\tau_{\rho\phi_i} \rho_{\pi_i} = \varphi_{\pi_i}(\Gamma, \pi_i)$ ;
        .
        .
        .
         $\tau_{\rho\phi_k} \rho_{\pi_k} = \varphi_{\pi_k}(\Gamma, \pi_k)$ 
        if (! $\rho_{\pi_1}$ .Equals( $\rho_{\pi_1}$ ))
            throw new TypeError("Type error");
        .
        .
    }
}
```

```

    .
    else if (! $\rho_{\pi_i}$ .Equals( $\rho'_{\pi_i}$ ))
        throw new TypeError("Type Error");
    .
    .
    else if (! $\rho_{\pi_k}$ .Equals( $\rho'_{\pi_k}$ ))
        throw new TypeError("Type Error");
    else
        return new  $\rho_{\sigma_j}$ ();
}
... //evaluation of other statements
}

public  $\sigma_j : \alpha$ 
{
     $\tau_{\sigma_{j_1}} \epsilon \sigma_{j_1}$ ;
    ...
     $\tau_{\sigma_{j_m}} \epsilon \sigma_{j_m}$ ;

    public  $\tau_{\rho_{\sigma_j}}$  Visit(Visitor< $\tau_{\rho_{\sigma_j}}$ > visitor)
    {
        return visitor. $\varphi_{\sigma_j}(\epsilon_{\sigma_{j_1}}, \dots, \epsilon_{\sigma_{j_m}})$ ;
    }
}

```

A general pseudo-code representation of the rule evaluation is shown in Algorithm 3.1.

Algorithm 3.1 Pseudocode of rule evaluation

```

function EVALUATE RULE( $R$  inference rule ,  $I$  input of the rule )
    if not  $R.Conclusion$  matches  $I$  then
        return error
    end if
    for all  $p$  in  $R.Premises$  do
         $p' \leftarrow$  textbf evaluate  $p$ 
        if not  $p.Result$  matches  $p$  then
            return error
        end if
    end for
    return  $R.Result$ 
end function

```

A graphical representation of the rule evaluation can be found in Figure 3.1.

3.1.2 Hard-coded implementation of Semantics

As shown in Section 2.5, there are multiple ways to express the semantics of a programming language. In this work we choose to make use of the operational semantics representation to have a uniform way of expressing both the type system and the semantics of a language. Let us consider again the semantics rule for **if-then-else** and **while-do** presented in Section 2.5. The operational semantics can be implemented generating the code in the object language that emulates the behaviour of the semantics rule, in

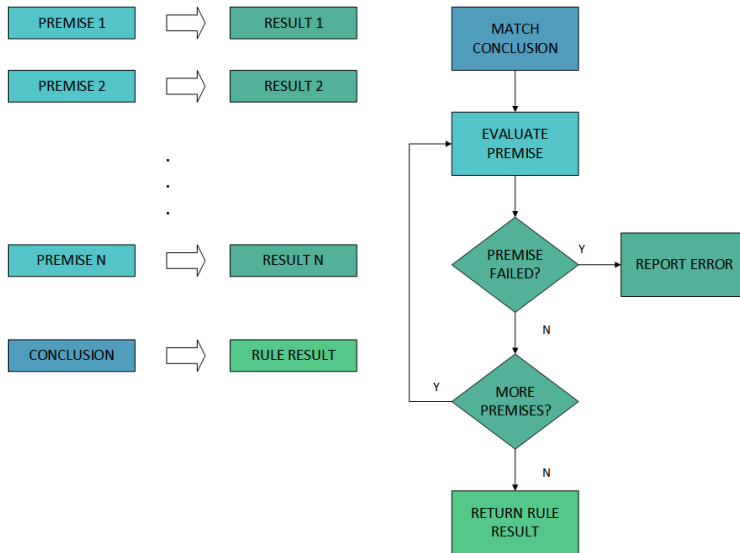


Figure 3.1: Diagram of rule evaluation: on the left side the structure of an inference rule, on the right side its evaluation expressed as a flow chart. The components of the rule are coloured to match the parts in which they are used in the flow chart.

the same fashion of a type rule. This process might first pass through an intermediate language, closer to the target language. In the case of an interpreter, the behaviour of the semantics must be implemented using the abstractions available in the host language. As an example, we show a possible implementation of the semantics of the two statements mentioned above in an interpreter both in a functional programming language and in an object-oriented language, as for the type rule.

For convenience, let us make a separate rule for the semantics of a sequence of statements from the specific semantics of the control structure. Also, we introduce the statement **skip** that performs no operation

$$\begin{array}{c}
 \frac{}{\langle \text{skip}; ks \rangle \Rightarrow \langle ks \rangle} \\
 \frac{\langle k \rangle \Rightarrow k'}{\langle k; ks \rangle \Rightarrow \langle k'; ks \rangle} \\
 \frac{\langle c \rangle \Rightarrow \text{true}}{\langle \text{if } c \text{ then } T \text{ else } E \rangle \Rightarrow T} \\
 \frac{\langle c \rangle \Rightarrow \text{false}}{\langle \text{if } c \text{ then } T \text{ else } E \rangle \Rightarrow E} \\
 \frac{\langle c \rangle \Rightarrow \text{true}}{\langle \text{while } c \text{ do } L \rangle \Rightarrow L ; \text{ while } c \text{ do } L} \\
 \frac{\langle c \rangle \Rightarrow \text{true}}{\langle \text{while } c \text{ do } L \rangle \Rightarrow \text{skip}}
 \end{array}$$

As for the type rules, we assume that the data type representing a statement is implemented by a discriminate union. The evaluation function first performs the pattern matching on the argument to select the correct rule to execute, in the same fashion of the type rule, but instead of analysing the types this time executes the specific behaviour of the statement, as specified by the semantics rule. For inference rules above we use the following code:

```

let rec interpretStmt (symbolTable : SymbolTable) (stmt : Statement) :
  Statement =
  match stmt with
  | Sequence(Skip,ks) -> interpretStmt symbolTable ks
  | Sequence(k,ks) ->
    let k' = interpretStmt symbolTable k
    interpretStmt symbolTable Sequence(k',ks)
  | If (cond,_then,_else) ->
    let condEvaluation = interpretExpr symbolTable cond
    if condEvaluation = True then
      _then
    else
      _else
  | While (cond,_do) ->
    let condEvaluation = interpretExpr symbolTable cond
    if condEvaluation = true then
      Sequence(_do,While(cond,_do))
    else
      Skip
  ... //other statements semantics

```

As for the type evaluation, we assume that `interpretExpr` is another function that is able to process expressions and return their value.

The function matches the kind of statement that we want to execute. In the case of a sequence of statements starting with a skip, we simply return the interpretation of the remaining part of the sequence (we indeed skip to the next statement), otherwise we have to run the first statement and then recursively evaluate the sequence formed by the result of the execution of the statement and the rest of the statements. This is needed, for instance, to correctly evaluate the body of a control structure. The body of each match case is responsible of emulating the intended behaviour described in the semantics of the control structure: the `if` returns the correct block to execute depending on the boolean value of the condition, instead the `while` returns either its body followed by the same `while` loop when the condition is `true`, otherwise `skip` to jump past the loop.

In the case of an object-oriented language, it is necessary to add a new implementation of the visitor pattern implementing the behaviour of the semantics for each statement:

```
public class StatementInterpreter : StatementVisitor
{
    ... //other statements semantics
    public Statement OnSequence(SymbolTable symbolTable, Sequence seq)
    {
        Statement k = seq.Head;
        Statement ks = seq.Tail;
        if (k.Equals(Skip))
            return ks.Visit(new StatementInterpreter());
        else
        {
            Statement k1 = k.Visit(new StatementInterpreter());
            Statement seq1 = new Sequence(k1, ks);
            return seq1.Visit(new StatementInterpreter());
        }
    }
    public Statement OnIf(Expression cond, Statement _then, Statement
        _else)
    {
        Value condValue = cond.Visit(new ExpressionInterpreter());
        if (condValue.Equals(new True()))
            return _then;
        else
            return _else;
    }
    public Statement OnWhile(Expression cond, Statement _do)
    {
        Value condValue = cond.Visit(new ExpressionInterpreter());
        if (condValue.Equals(new True()))
            return new Sequence(_do, new While(cond, _do));
        else
            return new Skip();
    }
    ... //other statement semantics
}
```

A further remark is that, for the sake of simplicity, here the interpretation only returns a new statement to execute obtained by processing the current statement, but in a real application it should also return a data structure representing the state of the program.

At this point it is possible to observe that this pattern can be generalized as well in a way analogous to that used for type rules for both implementa-

tions, which we omit for brevity.

3.1.3 Discussion

In Section 3.1.1 and 3.1.2 we have shown two implementations, one functional and one object-oriented, of type rules and semantics in a possible hard-coded compiler. We have also shown that the pattern can be generalized in both versions. Indeed, their behaviour must be hard-coded in the language chosen for the compiler implementation, regardless of the fact that the pattern is constantly repeated in every rule. This pattern can be captured in a meta-language that is able to process the type system and operational semantics definition of the language and generates the code in the target language necessary to execute the behaviour of the rules. In the following sections we describe the meta-language for *Metacasanova*, a meta-compiler that is able to read a program written in terms of type system/operational semantics rules defining a programming language, a program written in that language, and output executable code that mimics the behaviour of the semantics. The goal of this language is relieving the programmer from writing boiler-plate code when implementing a compiler for a (Domain-Specific) language.

3.2 Metacasanova overview

In this section we present the general idea behind Metacasanova. We start by defining the requirements of Metacasanova, then we proceed to give a general overview of the language, and finally we formalize the semantics of the language.

3.2.1 Requirements of Metacasanova

In order to relieve programmers of manually defining the behaviour described in Section 3.1.1 and 3.1.2 in the back-end of the compiler, we propose the following features for Metacasanova:

- It must be possible to define custom operators (or functions) and data containers. This is needed to define the syntactic structures of the language we are defining.
- It must be typed: each syntactic structure can be associated to a specific type in order to be able to detect meaningless terms (such as adding a string to an integer) and notify the error to the user.
- It must be possible to have polymorphic syntactical structures. This is useful to define equivalent “roles” in the language for the same syntactical structure; for instance we can say that an integer literal is both a *Value* and an *Arithmetic expression*.
- It must natively support the evaluation of semantics rules, as those shown above.

We can see that these specifications are compatible with the definition of meta-compiler, as the software takes as input a language definition written in the meta-language, a program for that language, and outputs runnable code that mimics the code that a hard-coded compiler would output.

3.2.2 Program structure

In this section we give an informal idea of how a Metacasanova program is organized. Further ahead this idea is expanded with additional details when we present the implementation details of the parser.

A Metacasanova program is mainly organized in three parts:

1. *Data and function declarations*: in this part it is possible to specify data structures, that define the syntactic constructs of the language, and functions used to evaluate terms of the language through rules.
2. *Type equivalence declarations*: in this part it is possible to specify polymorphism through type equivalence by stating that a type T_1 is equivalent to another type T_2 .
3. *Rule definitions*: in this part the programmer defines the type or semantics rules necessary to describe the type system or behaviour of the abstractions of the programming language.

A data structure or function declaration specifies the types of the arguments to construct the data structure or to pass to the function, their name, and the type of the data structure or the function itself

```
Data Expr -> "+" -> Expr : Expr
```

Note that Metacasanova allows you to specify any kind of notation for data types in the language syntax, depending on the order of definition of the argument types and the constructor name. In the previous example we used an infix notation. The equivalent prefix and postfix notations would be:

```
Data "+" -> Expr -> Expr : Expr
Data Expr -> Expr -> "+" : Expr
```

Optionally, it is possible to specify a precedence priority and the associativity. For example, the following code specifies that the multiplication has a higher precedence over the sum and that both are left-associative.

```
Data Expr -> "+" -> Expr : Expr Priority 0 <|
Data Expr -> "*" -> Expr : Expr Priority 1 <|
```

A function definition is similar to a data definition but it also has a return type. For instance the following is the evaluation function definition for the arithmetic expression above:

```
Func "eval" -> Expr : Value
```

A polymorphic data structure is defined through the keyword `is`, which specifies that a type T_1 is equivalent to another type T_2 . For example the following code specifies that a data structure of type `Value`, such as a list, can be used also as an expression of type `Expr`.


```
Data "$l" -> List : Value
Value is Expr
```

Metacasanova also allows to embed C# code into the language by using double angular brackets. This code can be used to embed .NET types when defining data or functions, or to run C# code in the rules. For example in the following snippet we define a floating point data which encapsulates a floating point number of .NET to be used for arithmetic computations:

```
Data "$f" -> <<float>> : Value
```

A rule in Metacasanova may contain a sequence of premises and a conclusion. The rule is executed if the input matches the pattern of the conclusion and all the premises return a result that matches the one specified in their rightmost part. In the following snippet we have the rule to evaluate the sum of two floating point numbers:

```
eval a => $f c
eval b => $f d
<<c + d>> => res
-----
eval (a + b) => $f res
```

Note that if one of the two expressions does not return a floating point value, then the entire rule evaluation fails. Also note that we can embed C# code to perform the actual arithmetic operation. Metacasanova selects a rule by means of pattern matching (in order of declaration of rules) on the function arguments. This means that both of the following rules will be valid candidates to evaluate the sum of two expressions:

```
...
-----
eval expr => res
...
-----
eval (a + b) => res
```

A more exhaustive explanation of the syntax of Metacasanova is given in Section 3.4, while an overview of the general shape of a program can be found in Figure 3.2.

3.2.3 Formalization

In what follows we assume that the pattern matching of the function arguments in a rule succeeds, otherwise a rule will fail to return a result. The informal semantics of the rule evaluation in Metacasanova is the following:

R1 A rule with no clauses or function calls always returns a result.

R2 A rule returns a result if all the clauses evaluate to **true** and all the function calls in the premise return a result.

R3 A rule fails if at least one clause evaluates to **false** or one of the function calls fails (returning no results).




Figure 3.2: Structure of a program in Metacasanova

We will express the semantics, as usual, in the form of logical rules, where the conclusion is obtained when all the premises are true. In what follows we consider a set of rules defined in the Metacasanova language R . Each rule has a set of function calls F and a set of clauses (boolean expressions) C . We use the notation f^r to express the application of the function f through the rule r . We will define the semantics by using the notation $\langle expr \rangle$ to mark the evaluation of an expression, for example $\langle f^r \rangle$ means evaluating the application of f through r . The following is the formal semantics of the rule evaluation in Metacasanova, based on the informal behaviour defined above:

$$\begin{aligned}
& \text{R1: } \frac{C = \emptyset \quad F = \emptyset}{\langle f^r \rangle \Rightarrow \{x\}} \\
& \text{R2: } \frac{\forall c_i \in C, \langle c_i \rangle \Rightarrow \text{true} \quad \forall f_j \in F, \exists r_k \in R \mid \langle f_j^{r_k} \rangle \Rightarrow \{x_{r,k}\}}{\langle f^r \rangle \Rightarrow \{x_r\}} \\
& \text{R3(A): } \frac{\exists c_i \in C \mid \langle c_i \rangle \Rightarrow \text{false}}{\langle f^r \rangle \Rightarrow \emptyset} \\
& \text{R3(B): } \frac{\forall r_k \in R, \exists f_j \in F \mid \langle f_j^{r_k} \rangle \Rightarrow \emptyset}{\langle f^r \rangle \Rightarrow \emptyset}
\end{aligned}$$

R1 says that, when both C and F are empty (we do not have any clauses or function calls), the rule in Metacasanova returns a result. R2 says that, if all the clauses in C evaluates to true and, for all the function calls in F we can find a rule that returns a result (all the function applications return a result for at least one rule of the program), then the current rule returns a result. R3(a) and R3(b) specify when a rule fails to return a result: this happens when at least one of the clauses in C evaluates to false, or when one of the function applications does not return a result for any of the rules defined in the program.

3.3 Architectural overview

In this section we provide a general overview of the architecture of Metacasanova compiler.

The compiler has a modular structure: in the front-end we find the the lexer/parser and a parser post-processing module. The latter is required because not all information necessary to build all the elements of the AST is immediately available during the parsing phase. For instance, some data structures store the file name that is being compiled and the name of the current module, but this information is available only after the parsing itself.

The generated AST is passed to the type checker to check the type correctness. Note that the type checker of the metacompiler checks the meta-types, i.e. the types defined in the meta-program, and not the types of the terms of the language that is being implemented in Metacasanova. This module checks the correctness of the declarations and the terms used in rules. The type checker outputs a data structure containing information about the types of the declarations and terms used in rules (in short a *typed program definition*).

The output of the AST is passed to the code generator that uses information about the types to correctly generate the target code. This is needed because Metacasanova generates C# code, which is a typed high-level language that requires information about the types to define variables, methods, and classes representing the elements of the meta-program.

Note that also, with this implementation choice, it is possible to support different high-level programming languages, both typed and untyped: the only component that changes will be the generation of the behaviour of the rules in the abstractions provided by the different target languages. A possible improvement of this architecture is generating a common intermediate language that is later translated into the target code, but this falls outside the scope of this work.

3.4 Parsing

In this section we explain in detail the grammar of Metacasanova and we present the architecture of its parser. The parser has been built in FsYacc (see Section 2.3.2) and completed by a post-processing module that executes some required transformation on the generated AST that are not convenient to perform during the parsing phase. As explained informally in Section 3.2, a Metacasanova program is made of four main sections: (i) a part containing inclusion directives, (ii) a part containing the declarations of the meta-data structures and evaluation functions used in the program, (iv) a part containing type equivalences, and (iii) a part containing evaluation rules that define the behaviour of the meta-program. The definition of the first part is trivial and we do not examine it in detail. We will instead describe in detail the other parts.

3.4.1 Declarations

Declarations contain *function* or *data* declarations. A meta-data structure is a meta-language representation of an abstraction of the language implemented in the metacompiler and contains both syntactical and structural information. For example, an arithmetic operator in a programming language can be represented as a meta-data structure containing both its symbol and the values of its arguments. Meta-data structures can be recursive, i.e. it is possible to have arguments that are instances of the same meta-data structure. This is done in order to allow recursive data structures such as lists. A meta-data structure declaration begins with the keyword **Data** and is followed by a series of arguments, which are separated by arrows, that can be both type names and strings representing the name of the meta-data. It is possible to declare an infix or suffix operator by placing its name after the first position of the arguments. For instance, the following code defines a sum operator for arithmetic expressions with an infix notation.

```
Data Expr -> "+" -> Expr : Expr
```

Type names are identifier that begins with an alphabetic character followed by one or more alphanumeric characters or underscores. The regular expression defining this syntax is expressed as

```
ID ::= ['a'-'z' 'A'-'Z'] ['a'-'z' 'A'-'Z' '_' '0'-'9']+
```

Type names can also contain external code enclosed by double angular brackets to make use of external types such as .NET primitive types (**int**, **float**, etc.). The operator names are strings that may contain any symbol usually allowed in strings in usual programming languages.

The arguments are followed by a type name defining the type of the meta-data structure. Optionally it is possible to specify a priority and the associativity, otherwise the default priority will be -1 and the operator will be left-associative.

Function declarations have the same structure except they begin with the keyword **Func** instead.

Both data and function declarations may define generic arguments. In order to specify generic arguments, they must be enclosed between square brackets after the declaration keyword. For instance the following code defines a data structure representing a tuple

```
Data[a,b] a -> "," -> b : Tuple[a,b]
```

Finally, each declaration must end with a line break. Line breaks are used in Metacasanova to separate different language elements, as in this case.

The grammar production used to describe the syntax of declarations is the following:

```
declaration:
| FUNC genericSeq typeOrNameDeclarations COLON typeDeclaration priority
  associativity newLineSeq {
Func(processParsedArgs $3 $5 (fst $1) (snd $1) $2 $6 $7) }
| DATA genericSeq typeOrNameDeclarations COLON typeDeclaration priority
  associativity newLineSeq {
Data(processParsedArgs $3 $5 (fst $1) (snd $1) $2 $6 $7) }
```

Since type names and data or function names can appear in any order, the parser generates a support polymorphic data structure in F#.

```
type TypeDeclOrName =
| Type of TypeDecl
| Name of string
```

This data structure is later transformed by the parser post-processor in a *symbol declaration*. A symbol declaration contains all the information about a declaration, including the data or function name, the types of the arguments, the priority, and the generic types. This information is later exploited by the type checker to verify the consistency of the declarations and type check the rest of the program.

Type equivalence declarations

Type equivalence has been presented as a separate part of the program but it has a tight relationship with the declarations. A type equivalence has the form

```
T1 is T2
```

where T1 and T2 are two meta-type names. They are used to specify that meta-type T1 is equivalent to meta-type T2 and can replace any argument of type T2 while constructing meta-data structures or calling functions. The grammar rule that defines a type equivalence declaration is

```
ID IS ID newLineSeq
```

where ID is the same regular expression used for type names. Again successive equivalence declarations should be separated by one or more line breaks. The grammar production generates a list of pairs in the AST containing the types involved in the equivalence. This data structure will be processed at a later stage by the type checker to generate an equivalence table.

3.4.2 Rules

Rules in Metacasanova are the language elements used to define the behaviour of the meta-program. A rule consists of a set of premises followed by a conclusion. Premises and conclusion are separated by a fraction line. Premises and conclusion are made of a left part consisting of a sequence of arguments, and a right part that can contain either a variable or a sequence of arguments. We call the left part of this syntactical structure *function call*, while the right part is the *result*. Premises differ from the conclusion as, besides function calls, they can also contain *bindings* and *clause*. Bindings are simply ways to rename values in the premises in the fashion of bindings in functional languages, while clauses are boolean predicates. Moreover, premises can also contain .NET code to directly emit, which can contain any C# code. The syntax of emitted code is the same as that of a normal premise, except the *function call* is replaced by the code to emit enclosed in double angular brackets. The following is the grammar production defining a premise¹:

```
premise:
| emit premises { $1 :: $2 }
| functionCall premises { $1 :: $2 }
| ID BIND arg newLineSeq premises { (Bind({ Namespace = ""; Name = fst
    $1 },Position.Create(snd $1,""),$3)) :: $5 }
| arg comparisonOp arg newLineSeq premises { (Conditional($1,$2,$3)) ::
    $5 }
| { [] }

functionCall:
| argSeq ARROW argSeq newLineSeq { FunctionCall($1,$3) }
```

Note that premises are optional (axioms do not have any premise in a logical rule), so an empty list is returned when none is given. Note that the namespace required for variables, such as in the binding, is left empty because at this point the namespace of the program is not available yet. The namespace will be later filled in by the parser post-processor. Also note that, in this stage, we do not check if each function call actually contains a function name, and we simply parse a premise (and a conclusion) as a sequence of arguments, that could be function or data names, variables, literals, or nested expressions. Nested expressions are expressions enclosed in brackets and are themselves other sequences of arguments. The actual control that premises and conclusions contain a function name is performed by the type checker, because to correctly identify the function name a complete symbol table, unavailable in this moment, is required. A conclusion has the same syntax of a function call:

```
conclusion:
| argSeq ARROW argSeq newLineSeq { ValueOutput($1,$3) }
```

¹ BIND is the symbol :=

The parser generates a data structure for a rule containing a representation of the premises and the conclusion:

- *function call*: A function call is simply a pair of list of arguments, where the left element is the call itself, while the right element is the result.
- *emit*: Emitted code contains the code in string format and the variable it is assigned to (used to save the result of expressions or function calls).
- *bind*: Bindings contain the variable name used for the binding and its argument, which can be a literal, the constructor of a meta-data structure, or another variable.
- *conditional*: Conditionals are boolean expressions that may contain comparison operators. Their representation stores their left and right argument and the comparison operator.

3.4.3 Parser post-processor

The parser post-processor is responsible to integrate in the AST all the information that is not directly available during the parsing phase. It is also responsible to re-arrange the terms appearing in the data and function declarations, and those appearing in function calls. Its main functions are the following: (i) insert the namespace and file information in the elements of the AST, (ii) rearrange the terms parsed from a declaration in a *symbol declaration* data structure, and (iii) parenthesize the function call terms according to their priority and associativity and rearrange them in a prefix notation.

The first task is trivial and will not be explained in details. Suffice to say that the process scans the AST starting from the root and recursively add the namespace and file name into all the nodes that must contain such information until a leaf is reached. Below we explain in detail the other how to accomplish the other two tasks.

Building the declaration data structure

As anticipated in Section 3.4.1, the name of the data or function and the types of its arguments can appear in any order. For convenience, the AST stores a data structure called *symbol declaration* that separates all the information about a declaration for further use, which is made of the following components:

- The name of the data or function.
- The type of the arguments of the data or function.
- The return type of the declaration: in the case of a data declaration this defines the type of the meta-data structure itself, while in a function declaration this defines the type of the result returned by the function.
- The operator arguments order, which can be prefix, infix, or suffix.
- The priority of the operator.
- The associativity of the operator.
- Possible generic arguments.

- The arity (amount of arguments) to the left and right of the operator symbol.

When the parser processes the arguments of a declaration, it creates a list of terms that can be either a type declaration or a name assigned to the meta-data or function in **string** format. The different arguments must be recognized and stored appropriately in the symbol declaration. In order to do so, the post-processor scans the arguments and, if the argument is a string, then it places it as a first element of a pair, otherwise it places the type declaration in a list of declarations. Algorithm 3.2 shows the details of this process. Note that it might be possible that the programmer commits the mistake of defining more than one name for the declaration, since we are still checking the syntax of the program. Thus the algorithm checks that the result of the function does not already contain a declaration name as, if it does, it means that a name argument has already been encountered while scanning the list.

The symbol declaration needs also to store the order of the declaration, its left and right arity, and the type of the declaration.

For the declaration order we have three options:

1. The first element of the parsed arguments is the declaration name. The declaration is then **prefix**.
2. The last element of the parsed arguments is the declaration name. The declaration is **suffix**.
3. If both 1 and 2 are false, then the name is in the middle and the declaration is **infix**.

Finding the left and right arity of an operator can be done simply by splitting the list of arguments in correspondence of the declaration name and then counting the elements of the two lists obtained by the split.

Finally, the post-processor must build the type of the declaration. Types in Metacasanova are represented in a way similar to typed lambda calculus [11, 15]: if a declaration as type arguments T_1, T_2, \dots, T_n , then its type representation is given as $T_1 \rightarrow T_2 \rightarrow \dots \rightarrow T_n$. This will allow at a later stage the type checking of partial function applications. The post-processor scans the arguments list and recursively add the element to a data structure representing an *arrow type*. The arrow type contains two elements corresponding to the elements to the left and right of the arrow. It is possible to build a chain of arrow types by recursively adding an arrow type as right argument of another arrow. For instance, to represent the type $A \rightarrow B \rightarrow C$ we use.

```
Arrow(A, Arrow(B, Arrow(C)))
```

The algorithm to build the type arrow simply scans the argument list and recursively add the current argument to the left of an arrow type and the result of the recursive call to the remaining part of the list as its right part. This process is shown in Algorithm 3.3. Note that in Metacasanova a declaration might contain no type arguments (only the name), thus we return an empty type as placeholder. If a declaration contains only one type argument then there is no need to build an arrow type and the type argument itself is returned. This is also used as a base case for the recursive build of an arrow type.

Algorithm 3.2 Arguments construction in a symbol declaration

```

function ARGSEPARATION(A list of arguments returned by the parser)
  name  $\leftarrow$  ""
  D  $\leftarrow \emptyset$ 
  for all a  $\in$  A do
    if a is a string then
      if name  $\neq$  "" then
        error: duplicate function name
      else
        name  $\leftarrow a$ 
      end if
    else
      D  $\leftarrow D \cup \{a\}$ 
    end if
  end for
  return (name, D)
end function

```

Algorithm 3.3 Type construction in a symbol declaration

```

function BUILDDECLARATIONTYPE(A list of arguments returned by the
parser)
  if A =  $\emptyset$  then
    return Empty
  else if A =  $\{x\}$  then
    return x
  else
    h  $\leftarrow$  HEAD(A)
    t  $\leftarrow$  TAIL(A)
    r  $\leftarrow$  BUILDDECLARATIONTYPE(t)
    return h  $\rightarrow$  r
  end if
end function

```

Parenthesization of function calls

As explained above, Metacasanova allows the declaration of functions and data types expressed with any notation (prefix, suffix, or infix) and with arbitrary associativity and precedence. Furthermore, this is possible regardless of the amount of arguments the data type or the function call use. Parsing operators according to precedence is a well-known problem that must be solved in order to avoid ambiguity in the language grammar. For instance, the expression $3 + 5 / 4$ can generate two different parse trees: $(3 + 5) / 4$ or $3 + (5 / 4)$. This ambiguity can be solved by setting the division operator to have a higher precedence over the sum operator, as in traditional arithmetic.

The first attempt to solve this problem was Dijkstra's Shunting-Yard algorithm [21] that takes an expression containing operators and a priority table and returns the same expression in reverse polish notation. This approach was later generalised by operator-precedence parsing which is available in all LALR(1) parser generators. Unfortunately, these approaches deal with binary operators, but are insufficient for operators of arbitrary arity. Moreover, parse generators such as Yacc allows to define a set of pre-defined language operators but do not allow to specify "custom" operators to extend the language with. A notable effort in parsing *mixfix operators* (i.e. operators with an arbitrary position in an expression) using precedence graphs has been done in [18].

In this work we used an AST-transformation technique that changes a function call expressed as a sequence of arguments into a parenthesized version based on defined priorities and associativity. A function call, as defined in Section 3.4.2, can be the left part of a premise or a conclusion, and is parsed as a sequence of arguments. Each argument can be represented in the following way:

- A literal.
- An identifier, that can start with an alphabetic character (*simple id*), or a symbol (such as %, #, &, @, ...) followed by a sequence of alphanumeric characters.
- A nested expression, which is any sequence of arguments enclosed between brackets (such as $(5 \text{ +++ } x)$).

We now give the definition of parenthesization of an argument sequence. In what follows we use the term *symbol* to indicate the name of *functions* or *meta-data* structure used in an argument sequence.

Definition 3.1. An argument sequence is parenthesized if all its arguments are (i) literals, (ii) identifiers, or (iii) a nested expressions containing a parenthesized argument sequence, and the nesting depth of a symbol is directly proportional to its priority (i.e. the highest-precedence operator is at maximum nesting depth).

For instance, given the precedence relation in Table 3.1, the following argument sequence is parenthesized

```
x -> ($ a1 (b1 % b2 b3) a2)
```

Listing 3.1: Example of parenthesization

Symbol	Priority	Left arity	Right arity	Associativity
%	2	1	2	Left
\$	1	0	3	Left
->	0	1	1	Left

Table 3.1: Precedence relation for Listing 3.1

The algorithm takes as input a precedence table with the structure of Table 3.1 and the argument sequence to parenthesize, and returns the parenthesized argument sequence. We adopt a recursive approach to the parenthesization problem, whose base case is that the sequence contains only identifiers and literals. Note that a sequence that contains a nested expression is in general not parenthesized because the sequence enclosed in brackets has not been parenthesized yet.

At this point we have two different possibilities for an argument sequence: (i) the sequence contains no symbol, or (ii) the sequence contains one or more symbols.

Parenthesization of sequences with no symbols When there are no symbols in the sequence of arguments, each argument can be a literal or identifier corresponding to no definition, or a nested expression. In the first case the argument is automatically parenthesized according to Definition 3.1. In the second case we must recursively parenthesize the sequence contained in the nested expression. If the result of this parenthesization is a single nested expression, then we take its content and store it into a single nested expression, to avoid redundant nesting parentheses of the form $((\dots(a_1 \ a_2, \dots, a_n)\dots))$. If the result is a series of arguments than we just place it inside a nested expression, obtaining something of the form $(a_1 \ a_2 \ (b_1 \ b_2 \ \dots) \ a_3 \ (c_1 \ c_2 \ \dots) \ \dots)$. It might be possible that the recursive algorithm tries to parenthesize a sequence containing no arguments (see the details of the algorithm below). In this case the algorithm simply returns an empty sequence.

Parenthesization of sequences with symbols Dealing with sequences containing symbols is more complex, since we must parenthesize them keeping into account the symbols priorities and associativity. According to Definition 3.1, in a parenthesized sequence the operator with the lowest priority must be at the top of the sequence nesting. The idea is then that at the current depth we should find the symbol with the lowest priority and a series of parenthesizations containing sequences with symbols of higher priority. The algorithm extracts the symbol in the sequence with the lowest priority and the positions in the sequence where it appears. Note that the same symbol might appear more than once. If the symbol associativity is left then we consider its rightmost occurrence in the sequence, otherwise its leftmost one. Now we split the sequence in two parts using this occurrence as separator and we recursively try to parenthesize these two parts. For instance, let us consider again the Precedence Table 3.1 and the following sequence of arguments

x -> \$ a1 b1 % b2 b3 a2

The algorithm will select the symbol \rightarrow as a separator, as it is the symbol with the lowest priority (there is only one occurrence, thus we neglect the part that selects the appropriate occurrence). It will then recursively parenthesize the sequences x and $\$ a_1 b_1 \% b_2 b_3 a_2$. At this point we have two parenthesizations of the left and right part: $\{l_1, l_2, \dots, l_n\}$ and $\{r_1, r_2, \dots, r_m\}$ respectively for the left and right sequence. Assuming that the left arity of the current symbol is a_l and the right one is a_r , then the parenthesization of the current operator will contain the elements $\{l_{n-a_l+1}, \dots, l_n\}$ and $\{r_1, \dots, r_{a_r}\}$. Assuming that the current symbol is σ , we now consider three cases:

1. The symbol uses a prefix notation: in this case the parenthesization will not contain any elements from $\{l_{n-a_l+1}, \dots, l_n\}$, thus the final parenthesization will be $\{l_1, l_2, \dots, l_n (\sigma, r_1, \dots, r_{a_r}) r_{a_r+1}, \dots, r_m\}$.
2. The symbol uses an infix notation: in this case the parenthesization will contain elements from both $\{l_{n-a_l+1}, \dots, l_n\}$ and $\{r_1, \dots, r_{a_r}\}$. The final parenthesization will be

$$l_1, \dots, l_{n-a_l} (l_{n-a_l+1}, \dots, l_n, \sigma, r_1, \dots, r_{a_r}) r_{a_r+1}, \dots, r_m$$

3. The symbol uses a suffix notation: in this case the parenthesization will not contain any elements from $\{r_1, \dots, r_{a_r}\}$ and the final parenthesization will be $l_1, \dots, l_{n-a_l} (l_{n-a_l+1}, \dots, l_n, \sigma) r_1, \dots, r_m$.

In order to clarify this process, let us consider again the argument sequence

$x \rightarrow \$ a_1 b_1 \% b_2 b_3 a_2$

and let us apply the algorithm to it (again using the Priority Table 3.1). The algorithm will test the whole sequence looking for symbols, and of course will find one. We then fall in the second part of the algorithm. The symbol with the lowest priority is \rightarrow , so the algorithm will recursively parenthesize x and $\$ a_1 b_1 \% b_2 b_3 a_2$. The left one is a base case of the recursion since the sequence contains only one identifier that is not a symbol, thus its parenthesization is the sequence itself. The right one contains other symbols so we have to recursively apply the algorithm. The operator with the lowest priority is now $\$$. The symbol is the first element of the sequence, thus the left subsequence obtained by the partitioning phase is empty (and the result of its parenthesization an empty sequence as well). The right subsequence is $a_1 b_1 \% b_2 b_3 a_2$. In this sequence there is only one symbol, which is $\%$, thus the algorithm will try to parenthesize $a_1 b_1$ and $b_2 b_3 a_2$. Their parenthesization is trivial and returns the sequences themselves. At this point we have to consider the arity of $\%$, which accepts one left argument and two right arguments. The algorithm will then enclose between brackets $b_1 \% b_2 b_3$. The full parenthesization leads then to $a_1 (b_1 \% b_2 b_3) a_2$. At this point this result is used to build the parenthesization of $\$$. This symbol accepts 3 right arguments (and no left argument), so the parenthesization will be $(\$ a_1 (b_1 \% b_2 b_3) a_2)$. Finally we have to use the result to build the parenthesization of \rightarrow . This symbol accepts 1 left argument and 1 right argument. The parenthesization of its left subsequence

was x , while its right parenthesization is $(\$ a1 (b1 \% b2 b3) a2)$, thus the final parenthesization will be $(x \rightarrow (\$ a1 (b1 \% b2 b3) a2))$. At this point, the outer parenthesization can be removed for better readability. The details of the algorithm are shown in Algorithm 3.4.

3.5 Type checking

The type checker of Metacasanova is responsible of two tasks: (i) checking that the declarations are correctly formed and (ii) check that the premises and conclusions of rules respect the meta-types defined in the declarations. We will now proceed to explain in detail how the two processes are implemented in the metacompiler.

3.5.1 Checking declarations

Checking declarations requires to check the consistency of *meta-type declarations* in each one of the function or data declarations. Note that, from now on, we will refer to meta-types simply as *types* for simplicity. Also we assume that, at this point, we have already built the *symbol table* for the meta-program containing all the symbol declarations with the complete information about the declaration itself. A type declaration in Metacasanova can have four different forms:

- *Zero*: A place-holder type used for function or meta-data that do not use any arguments.
- *External*: Used for embedded types from .NET.
- *Unsafe*: Unsafe type is a place-holder for external function calls, i.e. function defined in an external embedded language.
- *Argument*: A type argument is a simple type identifier followed by an optional list of generic arguments used for generic types. For example the type `Tuple[a,b]` is a type argument whose identifier is `Tuple` and whose generic arguments are a and b .
- *Arrow*: An arrow type has the form $T1 \rightarrow T2 \rightarrow \dots \rightarrow Tn$ and is used to represent the type of the arguments used when calling a function or when constructing a meta-data. For example the function declaration `Func Num -> "+" -> Num : Num` has the *Arrow* type `Num -> Num`. Note that the symbol declaration, for convenience, stores two different type declarations: the arguments types separated from the function returned type or the meta-data type and a *full type* that combines the type of the arguments and the returned type or data type into a single arrow. For example, for the function above, its full type would be `Num -> Num -> Num`.

The algorithm to check the type declarations behaves differently depending on the form of the type declaration. For *Zero* or *External* type declarations the check always succeeds. For *Arrow* the algorithm recursively check the left and right part of the arrow. In the case of an *Argument* we have two sub-cases: (i) the argument is an identifier with no generic arguments,

Algorithm 3.4 Parenthesization of a sequence of arguments. The operators $::$ and $@$ are respectively prepend and append on a list. With the notation $\langle S \rangle$ we denote a sequence S enclosed by parentheses.

```

function PARENTHESIZE( $S$  symbols in the sequence,  $A$  arguments sequence)
  if  $A = \emptyset$  then
    return  $A$ 
  else
    if  $s \neq \emptyset$  then
      let  $s'$  be the symbol with lowest priority in  $S$ 
      let  $I$  be the set of indices at which  $s'$  occurs in  $S$ .
       $l \leftarrow \emptyset$ 
       $r \leftarrow \emptyset$ 
      if  $s'$  is left-associative then
         $I' \leftarrow \text{LAST}(I)$ 
         $l, r \leftarrow \text{SPLITAT}(I')$ 
         $r \leftarrow \text{TAIL}(r)$ 
      else
         $I' \leftarrow \text{FIRST}(I)$ 
         $l, r \leftarrow \text{SPLITAT}(I')$ 
         $r \leftarrow \text{TAIL}(r)$ 
      end if
      let  $lsym$  be the symbols in  $l$ 
      let  $rsym$  be the symbols in  $r$ 
       $lpar \leftarrow \text{PARENTHESIZE}(lsym, l)$ 
       $rpar \leftarrow \text{PARENTHESIZE}(rsym, r)$ 
      let  $larity$  be left arity of  $s'$ 
      let  $rarity$  be right arity of  $s'$ 
       $largs, plargs \leftarrow \text{SPLITAT}(|largs| - larity)$ 
       $rargs, prargs \leftarrow \text{SPLITAT}(rarity)$ 
      if  $larity + rarity > 0$  then
        if  $s'$  is prefix then
           $e \leftarrow s' :: prargs$ 
        else if  $s'$  is suffix then
           $e \leftarrow plargs @ s'$ 
        elses  $s'$  is infix
           $e \leftarrow plargs @ s' @ prargs$ 
        end if
        return  $largs @ \langle e \rangle @ rargs$ 
      else
        return  $largs @ s' @ rargs$ 
      end if
    else
       $p \leftarrow \emptyset$ 
      for all  $a \in A$  do
        if  $a = \langle e \rangle$  then
          let  $S'$  be symbols in  $e$ 
           $p' \leftarrow \text{PARENTHESIZE}(S') e$ 
          if  $p' = \langle e' \rangle$  then
             $p \leftarrow p @ e'$ 
          else
             $p \leftarrow p @ \langle p' \rangle$ 
          end if
        else
           $p \leftarrow p @ \{a\}$ 
        end if
      end for
    end if
  end if
end function

```

or (ii) the argument is an identifier followed by a number of generic arguments. The first case is simple, as it is enough to check whether the type is defined in the symbol table or not. If the type cannot be found in the symbol table then it is undefined and an error is returned. In the case of a type with generic arguments we must check that the amount of provided generics matches the amount required for the generic type; then we must check if the provided generic arguments have a correct form. A generic argument can be an identifier or again a type accepting other generic arguments. This is the case, for instance, of a type declaration such as `Tuple[List[int], Tuple[a, List[float]]]`. In the case of a simple generic identifier we must only check that the generic identifier is in the scope of the declaration. For instance the declaration

```
Func[a,b] "foo" -> a -> b : b
```

would be a valid declaration since the generic identifier are in the scope of the declaration, while

```
Func[a,b] "foo" -> a -> b : c
```

would be invalid if `c` is not a data type defined in the meta-program. In the case of nested types the algorithm must recursively check again that the type exists and that the generic arguments are valid. The procedure details are described in Algorithm 3.5.

Algorithm 3.5 Type checking of a symbol declaration

```

function CHECKDECLARATION( $S$  symbol table,  $G$  declared generic arguments,  $d$  type declaration)
  if  $d = \text{Empty}$  or  $d = \text{Zero}$  then
    return
  else
    let  $G'$  be the generics required for  $d$ 
    if  $G' \neq \emptyset$  then
      if  $|G'| \neq |G|$  then
        error: invalid amount of generic arguments
      else
        for all  $g \in G'$  do
          CHECKDECLARATION( $S, G, g$ )
        end for
      end if
    else
      if  $d \in S$  or  $d \in G$  then
        return
      else
        error: undefined type
      end if
    end if
  end if
end function

```

3.5.2 Checking rules

Type checking a rule requires to type check its conclusion and all of the premises. In what follows we assume that the parser post-processor has already parenthesized and normalized all the function calls, so that every function call is parenthesized according to symbol priority and associativity and that the symbol name is in the first position of an argument sequence. Moreover, we use the following definition relative to meta-data arguments:

Definition 3.2. An argument is said to be an *explicit data argument* when it is an expression constructing a meta-data.

For example, in the following meta-program

```
Data Expr -> "+" -> Expr : Expr
Func "eval" -> Expr -> Environment : Value
...

eval a env -> a'
eval b env -> b'
<<a' + b'>> -> v
-----
eval (a + b) env -> v
```

Listing 3.2: Example of an explicit data argument in Metacasanova

the first argument of `eval` in the conclusion is an explicit data argument. Moreover we use the following definition to define the compatibility of two types.

Definition 3.3. Let T be the set of types defined in a meta-program and $E = \{(t_i, t_j) \mid t_1 \in T \wedge t_2 \in T\}$, then we say that the type t_1 is *compatible* with t_2 if either $t_1 = t_2$ or $\exists(t_i, t_j) \in E \mid t_1 = t_i \wedge t_2 = t_j$.

Checking the conclusion

A conclusion must always contain a function call. In order to type-check the function call correctly, we must check that (i) the arity of the function is respected, i.e. that the arguments are not more than what the function expects (the type system of Metacasanova supports partial function application, so it is allowed to pass less arguments), and (ii) that the type of each argument is compatible with what provided in the declaration. Moreover, since a conclusion might contain explicit data arguments, we have to recursively add all the variables contained in the explicit data arguments to the local variables of the current function, as they could be used in the premises. As an example, refer to Listing 3.2, where the variables `a` and `b` are defined in the argument of `eval` and later used in the premises. Checking *i* is trivial: it is sufficient to compare the length of the given argument sequence with the length of the arguments provided in the declaration. If the length of the argument sequence is greater than the arguments defined in the declaration then an error is returned. Checking *ii* is more complex and the details are explained below.

The right hand-side of a conclusion might contain a variable or an explicit data argument. In both cases its type checking must be delayed until all premises are processed because a variable appearing in one of them might be used. For example, in Listing 3.2, variable `v` is used in the right hand-side of the conclusion and defined in the result of the last premise.

Checking a premise

A premise, as explained in Section 3.4.2, can be (i) a function call, (ii) a binding, or (iii) a clause.

In case *i* we have to type check the arguments of the function call in the same fashion of the conclusion but with a slight difference: when encountering a variable this must not be added to the local variable set but rather looked up in it. If the lookup fails then the variable is undefined and cannot be used. The same happens when checking the arguments of explicit data arguments. If the call is correctly typed, then we must check its result. The result of a call can be either a variable or an explicit data argument. In the first case the variable is added to the local variables and its type set to the return type of the function. In the case of an explicit data argument then all the arguments that are variable are added to the local variables with the appropriate type read from the meta-data structure declaration and, in case of a nested explicit data argument, the procedure is recursively applied.

In case *ii* the binding is correctly typed if its right argument is correctly type. The right argument can be a variable or an explicit data argument so we apply the same method that we used to check variables and explicit data arguments in function calls. If this test succeeds then the left argument of the binding, which is always a variable, is added to the local variables with the type of the right argument.

In case *iii* the clause is correctly typed if the types of the arguments used in the comparison operator are compatible with respect to the operator itself and if they are themselves correctly typed. For example, for the equality comparison, we must ensure that both arguments are correctly typed. As always, if external types are involved the test automatically succeeds because nothing can be known at this point about the compatibility of the provided types. Of course the test fails if we use a comparison operator combining external types with types defined in the meta-program, because they will always be incompatible.

Checking a single argument

When checking the type of an argument we have to consider several cases depending on what kind of argument we are inspecting. The reader can find the details of each case below. Note that we will never consider external types as their type checking is delegated to the external code compiler that is used when compile the generated code, so their type checking in the Metacasanova type checker is always successful.

Checking literals If the argument is a literal then we have to consider three sub-cases:

1. The expected type is generic. In this case, since we are providing explicitly a literal, the generic can be assigned the specific type of the literal.

2. The expected type is non-generic. In this case we have to check if the type of the literal is compatible with the expected type. This only happens if the expected type is one of the native types supported by Metacasanova.
3. The expected type is a meta-data structure requiring generic arguments. In this case the type is always incorrect, since a literal is always incompatible with meta-data.

Checking identifiers When we have an identifier as argument, this might be either a symbol for a meta-data structure taking no arguments, or simply a variable. In the first case we simply check that the type of the meta-data structure is compatible with the expected type. In the second case the result depends on whether we are type checking a conclusion or a premise. If we are checking a conclusion, then the variable must be added to the local variables in the scope of the rule, otherwise we must look up the local variables to check if it was previously defined. If the lookup fails then the variable is undefined and an error is returned.

Checking nested expressions Checking a nested expression requires, in the first place, to recursively type check the arguments used in the nested expression. If this check succeeds then we must also ensure that the type of the meta-data structure that we are analysing is compatible with the expected type.

Checking the type compatibility According to Definition ??, a type is compatible with another if they are equal or if, in the symbol table, there exists a specified equivalence between the first type and the second. We have thus to distinguish two cases: (i) check if the types are equal and, if this fails, (ii) check if the provided type is paired with the expected type in an equivalence table. For the following considerations refer to the type structure defined in Section 3.5.1.

Testing type equality Checking type must consider three options: (i) the type is a simple identifier, (ii) the type is an Arrow type, and (iii) the type is External, Unsafe or Zero.

In case *i* we must compare two Type Arguments t_1 and t_2 other wise the test fails immediately. In this case it is enough to simply check that $t_1 = t_2$.

In case *ii* we must compare two Arrow types otherwise the test fails immediately. Let us assume that we have $T ::= t_1 \rightarrow t_2 \rightarrow \dots \rightarrow t_n$ and $U ::= u_1 \rightarrow u_2 \rightarrow \dots \rightarrow u_m$, then we check if $t_1 = t_2$ and recursively apply the type equality test on $t_2 \rightarrow \dots \rightarrow t_n$ and $u_m \rightarrow \dots \rightarrow u_m$. Note that if $n \neq m$ at some point the test will fail because we will compare a Type Argument with an Arrow Type, which will always fail.

In case *iii* the test succeeds if one of the types (either the provided one or the expected one) is external or unsafe; it also succeeds if both types are Zero types.

Testing type compatibility Type equivalence might succeed in only two cases: (i) when testing two Argument Types, and when testing two Arrow Types; in all other cases the test fails immediately. In what follows we write $t_1 \equiv t_2$ to say that t_1 is compatible with t_2 .

In case *i* we have to check if the provided type is paired with the expected type in a type equivalence map stored in the symbol table. If the lookup in the map is unsuccessful then an error is returned.

In case *ii* again we consider $T ::= t_1 \rightarrow t_2 \rightarrow \dots \rightarrow t_n$ and $U ::= u_1 \rightarrow u_2 \rightarrow \dots \rightarrow u_m$. This time we check if $t_1 \equiv u_1$ and then we recursively check the equivalence of $t_2 \rightarrow \dots \rightarrow t_n$ with $u_2 \rightarrow \dots \rightarrow u_m$. Again, the equivalence test fails if $n \neq m$ because we will end up comparing an Argument Type with an Arrow Type.

3.6 Code generation

Metacasanova uses C# as target language and generates code compatible with any library compiled in the .NET framework. In this phase we must (i) generate the appropriate abstractions in C# to represent meta-data structures and their type equivalence, and (ii) generate the code to implement the semantics of the rule evaluation, as described in Section 3.2.3.

3.6.1 Meta-data structures code generation

The type of each data structure is generated as an interface in C#. Each data structure defined in Metacasanova is mapped to a `class` in C# that implements such interface. The class contains as many fields as the number of arguments the data structure contains. Each field is given an automatic name `argC` where `C` is the index of the argument in the data structure definition. The data structure symbols used in the definition might be pre-processed and replaced in order to avoid illegal characters in the C# class definition. The class contains an additional field that stores the original name of the data structure before the replacement is performed, used for its “pretty print”. For example the data structure.

```
Data "$i" -> int : Value
```

will be generated as

```
public interface Value { }

public class __opDollari : Value
{
    public string __name = "$i";
    public int __arg0;

    5
    public override string ToString()
    {
        return "(" + __name + " " + __arg0 + ")";
    }
}
```

3.6.2 Code generation for rules

Each rule contains a set of premises that in general call different functions to produce a result, and a conclusion that contains the function evaluated by the current rule and the result it produces. The code generation for the rules follows the steps below:

1. Generate a data structure for each function defined in the meta-program.
2. For each function f extract all the rules whose conclusion contains f .
3. Create a **switch** statement with a case for each rule that is able to execute the function (the function is in its conclusion).
4. In the case block of each rule, define the local variables defined in the rule.
5. Apply pattern matching to the arguments of the function contained in the conclusion of the rule. If it fails, jump immediately to the next case (rule).
6. Store the values passed to the function call into the appropriate local variables.
7. Run each premise by instantiating the class for the function used by it and copying the values into the input arguments.
8. Check if the premise outputs a result and, in the case of an explicit data structure argument, check the pattern matching. If the premise result is empty or the pattern matching fails for all the possible executions of the premise then jump to the next case.
9. Generate the result for the current rule execution.

In what follows, we use as an example the code generation for the following rule (which computes the sum of two integer expressions in a programming language):

```
eval a -> $i c
eval b -> $i d
<< c + d >> -> e
-----
eval (a + b) -> $i e
```

From now on we will refer to an argument as *explicit data argument* when its structure appears explicitly in the conclusion or in one of the premises, as in the case of $a + b$ in the example above.

Data structure for the function

As first step the meta-compiler generates a class for each function defined in the meta-program. This class contains one field for each argument the function accepts. It also contains a field to store the possible result of its evaluation. This field is a **struct** generated by the meta-compiler defined as follows:

```
public struct __MetaCnvResult<T> { public T Value; public bool HasValue;
}
```

The result contains a boolean to mark if the rule actually returned a result or failed, and a value which contains the result in case of success.

For example, the function

```
Func eval -> Expr : Value
```

will be generated as

```
public class eval
{
    public Expr __arg0;
    public __MetaCnvResult<Value> __res;
    ...
}
```

Rule execution

The class defines a method **Run** that performs the actual code execution. The meta-compiler retrieves all the rules whose conclusion contains a call to the current function, which define all the possible ways the function can be evaluated with. It then creates a **switch** structure where each **case** represents each rule that might execute that function. The result of the rule is also initialized here (the **struct** will contain a default value and the boolean flag will be set to **false**). Each **case** defines a set of local variables, that are the variables used within the scope of that rule.

Local variables definitions and pattern matching of the conclusion

At the beginning of each **case**, the meta-compiler defines the local variables initialized with their respective default values. It also generates then the code necessary for the pattern-matching of the conclusion arguments. Since variables always pass the pattern-matching, the code is generated only for arguments explicitly defining a data structure (see the examples about arithmetic operators in Section ??) and literals. If the pattern matching fails then the execution jumps to the next **case** (rule). For instance, the code for the following conclusion

```
...
-----
eval (a + b) -> $i e
```

is generated as follows

```

case 0:
{
  Expr a = default(Expr);
  Expr b = default(Expr);
  int c = default(int);
  int d = default(int);
  int e = default(int);
  if (!(__arg0 is __opPlus)) goto case 1;
  ...
}

```

Note that an explicit data argument, such in the example above, might contain other nested explicit data arguments, so the pattern-matching is recursively performed on the data structure arguments themselves.

Copying the input values into the local variables

When each function is called by a premise, the local values are stored into the class fields of the function defined in Section 3.6.2. These values must be copied to the local variables defined in the **case** block representing the rule. Particular care must be taken when one argument is an explicit data. In that case, we must copy, one by one, the content of the data into the local variables bound in the pattern matching. For example, in the rule above, we must separately copy the content of the first and second parameter of the explicit data argument into the local variables **a** and **b**. The generated code for this step, applied to the example above, will be:

```

__opPlus __tmp0 = (__opPlus).__arg0;
a = __tmp0.__arg0;
b = __tmp0.__arg1;

```

Note that the type conversion from the polymorphic type **Expr** into **opPlus** is now safe because we have already checked during the pattern matching that we actually have **opPlus**.

Generation of premises

Before evaluating each premise, we must instantiate the class for the function that they are invoking. The input arguments of the function call must be copied into the fields of the instantiated object. If one of the arguments is an explicit data argument, then it must be instantiated and its arguments should be initialized, and then the whole data argument must be assigned to the respective function field. After this step, it is possible to invoke the **Run** method of the function to start its execution. The first premise of the example above then becomes (the generation of the second is analogous):

```

eval a -> $i c

```

```

eval __tmp1 = new eval();
__tmp1.__arg0 = a;
__tmp1.Run();

```

Checking the premise result

After the execution of the function called by a premise, we must check if a rule was able to correctly evaluate it. In order to do so, we must check that the result field of the function object contains a value, and if not the rule fails and we jump to the next case (rule), which is performed in the following way:

```
if (!(__tmp1.__res.HasValue)) goto case 1;
```

If the premise was successfully evaluated by one rule, then we must check the structure of the result, which leads to the following three situations:

1. The result is bound to a variable.
2. The result is constrained to be a literal.
3. The result is an explicit data argument.

In the first case, as already explained above, the pattern matching always succeeds, so no check is needed. In the second case, it is enough to check the value of the literal. In the last case, all the arguments of the data argument must be checked to see if they match the expected result. In general this process is recursive, as the arguments could be themselves other explicit data arguments. If the result passes the check, then the result is copied into the local variables, in a fashion similar to the one performed for the function premise. For instance, for the premise

```
eval a -> $i c
```

the meta-compiler generates the following code to check the result

```
if (!(__tmp1.__res.Value is __opDollari)) goto case 1;
__MetaCnvResult<Value> __tmp2 = __tmp1.__res;
__opDollari __tmp3 = (__opDollari)__tmp2.Value;
c = __tmp3.__arg0;
```

Generation of the result

When all premises correctly output the expected result, the rule can output the final result. In order to do that, the generated code must copy the right part of the conclusion (the result) into the **res** variable of the function class. If the right part of the conclusion is, again, an explicit data argument, then the data object must first be instantiated and then copied into the result. For example the result of the rule above is generated as follows:

```
res = c + d;
__opDollari __tmp7 = new __opDollari();
__tmp7.__arg0 = res;
__res.HasValue = true;
__res.Value = __tmp7;
break;
```

After this step, the rule evaluation successfully returns a result.

This implementation choice is due to the fact that we plan to support partial function applications, thus, when a function is partially applied, there

is the need to store the values of the arguments that were partially given. This could still be implemented with static methods and lambdas in C#, but not all programming languages natively support lambda abstractions, so we chose to have a set-up that allows us to change the target language without dramatically altering the logic of code generation.

Chapter 4

Language design in Metacasanova

Two languages implementing
the same idea must, on pain of
death, use different terms

Feldman's Law of
Programming

4.1 The C-- language

4.2 Casanova 2.5 in Metacasanova

In this section we will briefly introduce the Casanova language, a domain specific language for games. We then show a re-implementation, which we call Casanova 2, of the Casanova 2 language hard-coded compiler as an example of use of Metacasanova.

4.2.1 The Casanova language

Casanova 2.5 is a language oriented to video game development which is based on Casanova 2 [4]. A program in Casanova is a tree of *entities*, where the root is marked in a special way and called *world*. Each entity is similar to a *class* in an object-oriented programming language: it has a constructor and some fields. The fields do not have access modifiers because they are not directly modifiable from the code except with a specific statement. Each entity also contains a list of *rules*, that are methods that are ticked in order with a specific refresh rate called *dt*. Each rule takes as input four elements: *dt*, *this*, which is a reference to the current entity, *world* that is a reference to the world entity, and a subset of entity fields called *domain*. A rule can only modify the fields contained in the domain. The rules can be paused for a certain amount of seconds or until a condition is met by using the *wait* statement. It is possible to modify the values of the fields in the domain by

using the `yield` statement which takes as input a tuple of values to assign to the fields. When the `yield` statement is executed the rule is paused until the next frame. Also the body of control structures (`if-then-else`, `while`, `for`) is interruptible. In the following section we show the implementation of Casanova 2.5 in Metacasanova.

4.2.2 Casanova 2.5

The memory in Casanova 2.5 is represented using three maps, where the key is the variable/field name, and the value is the value stored in the variable/field. The first dictionary represents the global memory (the fields of the `world` entity or *Game State*), the second dictionary represents the current entity fields, and the third the variable bindings local to each rule.

The core of the entity update is the `tick` function. This function evaluates in order each rule in the entity by calling the `evalRule` function. This function executes the body of the rule and returns a result depending on the set of statements that has been evaluated. This result is used by `tick` to update the memory and rebuild the rule body to be evaluated at the next frame. The result of `tick` is a `State` containing the rules updated so far, and the updated entity and global fields. Since a rule must be restarted after the whole body has been evaluated, we need to store a list containing the original rules, which will be restored when evaluation returns `Done` (see below). At each step the function recursively calls itself by passing the remaining part of original rules (the rules which body was not altered by the evaluation of the statements) and modified rules (which body has been altered by the evaluation of the statements) to be evaluated. The function stops when all the rules have been evaluated, and this happens when both the original and the modified rule lists are empty.

Interruption is achieved by using *Continuation passing style*: the execution of a sequence of statements is seen as a sequence of steps that returns the result of the execution and the remaining code to be executed. Every time a statement is executed we rebuild a new rule whose body contains the continuation which will be evaluated next.

The possible results returned by the `tick` function are the following: (i) `Suspend` contains a `wait` statement with the updated timer, the continuation, and a data structure called `Context` which contains the updated local variables, the entity fields, and the global fields. The function rebuilds a rule which body is the sequence of statements contained by the `Suspend` data structure. (ii) `Resume` is returned when the timer must resume after the last waited frame. In order not to skip a frame we must still re-evaluate the rule at the next frame and not immediately. In this case the argument of `Resume` is only the remaining statements to be executed. (iii) `Yield` stops evaluation for one frame. We use the continuation to rebuild the rule body. Memory is updated by `evalRule`. (iv) `Done` stops the evaluation for one frame and rebuilds the original rule body by taking it from the original rules list.

For brevity we write only the code for `Suspend`. A full implementation can be found at [25]. You can see a schematic representation of the tick function in Figure 4.1.

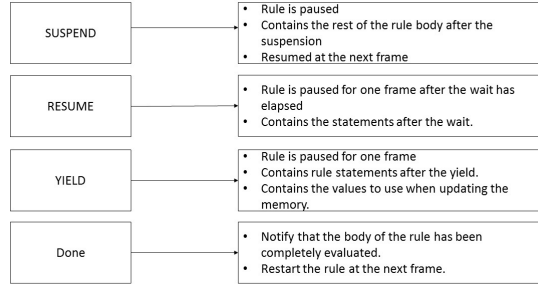


Figure 4.1: Casanova 2.5 rule evaluation

```

evalRule (rule dom body k locals delta) fields globals => Suspend (s;
    cont) (Context newLocals newFields newGlobals)
r := rule dom s cont newLocals dt
tick originals rs newFields newGlobals dt => State updatedRules
    updatedFields updatedGlobals
st := State (r::updatedRules) updatedFields updatedGlobals
-----
tick (original::originals) ((rule dom body k locals delta)::rs) fields
    globals dt => st

```

The function `evalRule` calls `evalStatement` to evaluate the first statement in the body of the rule passed as argument. The result of the evaluation of the statement is processed in the following way: (i) if the result is `Done`, `Suspend` or `Resume` then it is just returned to the caller function. We omit the code for this case, since it is trivial; (ii) if the result is `Atomic` it means that the evaluated statement was uninterruptible and the remaining statements of the rule must be re-evaluated immediately; (iii) if the result is `Yield` then the fields in the domain are updated recursively in order and then the updated memory is encapsulated in the `Yield` data structure and passed to the caller function.

```

evalStatement b k ctxt dt => Atomic z c
evalRule (rule dom z nop c dt) => res
-----
evalRule (rule dom b k ctxt dt) => res

```

```

evalStatement b k (Context locals fields globals) dt => Yield ks values
    context
updateFields dom values context => updatedContext
-----
evalRule (rule dom b k locals dt) fields globals => Yield ks values
    updatedContext

```

Note that, in case of a rule containing only atomic statements, we will eventually return `Done` after having recursively called `evalStatement` for all the statements, and the rule will be paused for one frame. The `evalStatement` function is used both to evaluate a single statement and a sequence of statements. When evaluating a sequence of statements,

the first one is extracted. A continuation is built with the following statement and passed to a recursive call to `evalStatement` which evaluates the extracted statement. If the existing continuation is non-empty, then it is added before the current continuation. If both the continuation and the body are empty (situation represented by the `nop` operator) then it means the rule evaluation has been completed and we return `Done`.

```

a != nop
-----
addStmt a b => a;b                                addStmt nop nop => nop

addStmt b k => cont
evalStatement a cont ctxt dt => res
-----
evalStatement (a;b) k ctxt dt => res                evalStatement nop nop
      ctxt dt => Done ctxt

```

We will now present, for brevity, only the evaluation of the `wait` and `yield` statements. Both the evaluation of the control structures and the variable bindings always return `Atomic` because they do not, by definition, pause the execution of the rule.

The `wait` statement has two different evaluations, based on the rules defined in Section ??: (i) the timer has elapsed: in this case we return `Resume` which contains the code to execute after the `wait` statement, or (ii) the timer has not elapsed: in this case we return `Suspend` which contains the `wait` statement with the updated timer followed by the continuation.

```

<<t <= dt>> == false
-----
evalStatement (wait t) k ctxt dt => Suspend wait <<t - dt>>;k ctxt

<<t <= dt>> == true
-----
evalStatement (wait t) k ctxt dt => Resume k ctxt

```

The `yield` statement takes as argument a list of expressions whose values are used to update the corresponding fields in the rule domain. The evaluation rule recursively evaluates the expressions and stores them into a list passed as argument of the `Yield` result. Those arguments are used later by `evalRule` to update the corresponding fields.

```

eval expr ctxt => v
evalYield exprs ctxt => vs
-----
evalYield (expr :: exprs) ctxt => v :: vs                evalYield nil ctxt
=> nil

```

In this section we provide an implementation of a patrol script for an entity in a game. The sample is made up of an entity, representing a guard, and a couple of checkpoints. The guard continuously moves between the two checkpoints. We choose this sample because this is a typical behaviour implemented in several games, where the user is able to set up a patrol route for a unit. We show the comparison between the sample implemented in Casanova 2.5 and an equivalent implementation in Python with respect to the running time. We then show a comparison between the hard-coded compiler of Casanova 2.0 and the implementation of Casanova 2.5 in Metacasanova with respect to the code length.

4.2.3 Chosen languages

We compared the running time of the sample in metacompiled Casanova with an equivalent implementation in Python. This language was chosen based on its use in game development: Python has been used extensively in several games such as Civilization IV [23] or World in Conflict [32] because of the native support for coroutines. We deliberately ignore C++ and C# implementations, although they are widely used in the industry, because we knew in advance [4] that the current version of the code generated by the meta-compiler would not match the high performance of these languages: the main goal of this work is to reduce the effort of writing a compiler for a DSL for games while having acceptable performance.

4.2.4 Performance

The performance results are shown in Table 4.1. We see that the generated code has performance on the same order as Python. This is mainly due to the fact that the memory, in the metacompiled implementation of Casanova, is managed through a map, and because of the virtuality of the implemented operators. Each time Casanova accesses a field in an entity this must be looked up into the map. To this we add the complexity of dynamic lookups when we must deal with polymorphic results into the rules.

From Table 4.2 we see that the implementation of Casanova 2.0 language in Metacasanova is almost 5 times shorter in terms of lines of code than the previous Casanova implementation in F#. We believe it is worthy noticing that structures with complex behaviours, such as *wait* or *when*, require hundreds of lines of codes with a standard approach (the code lines to define the behaviour of the structure plus the support code to correctly generate the state machine), while in the meta-compiler we just need tens of lines of codes to implement the same behaviour. Moreover we want to point out that the previous Casanova compiler was written in a functional programming language: these languages tend to be more synthetic than imperative languages, so the difference with the same compiler implemented in languages such as C/C++ might be even greater.

The readability with respect to the hard-coded compiler code is also improved: we managed to implement the behaviour of synchronization and timing primitives almost imitating one to one the formal semantics of the language definition (see the semantics rules in Section ?? and their implementation in Section ??). In the hard-coded compiler implementation for Casanova 2.0 the semantics are lost in the code for generating finite state machines.

4.2.5 Discussion

Metacasanova has been evaluated in [19] by re-building the DSL for game development Casanova [5, 3]. Even though the size of the code required to implement the language has been drastically reduced (almost 1/5 shorter), performance dropped dramatically. We identified a main problem causing the performance decay that, if solved, will improve the performance of the generated code.

Casanova 2.5		
Entity #	Average update time (ms)	Frame rate
100	0.00349	286.53
250	0.00911	109.77
500	0.01716	58.275
750	0.02597	38.506
1000	0.03527	28.353
Python		
Entity #	Average update time (ms)	Frame rate
100	0.00132	756.37
250	0.00342	292.05
500	0.00678	147.54
750	0.01087	91.988
1000	0.01408	71.002

Table 4.1: Patrol sample evaluation

Casanova 2.5 with Metacasanova	
Module	Code lines
Data structures and function definitions	40
Query Evaluation	16
While loop	4
For loop	5
If-then-else	4
When	4
Wait	6
Yield	10
Additional rules for Casanova program evaluation	40
Additional rules for basic expression evaluation	201
Total:	300
Casanova 2.0 compiler	
Module	Code lines
While loop	10
For-loop and query evaluation	44
If-Then-Else	15
When	11
Wait	24
Yield	29
Additional structures for rule evaluation	63
Structures for state machine generations	754
Code generation	530
Total:	1480

Table 4.2: meta-compiler vs standard compiler

In order to encode a symbol table in the meta-compiler in the current implementation (used for example to store the variables defined in the local scope of a control structure or to model a class/record data structure), we are left with two options: (i) define a custom data structure made of a list of pairs, containing the field/variable name as a string and its value, in the following way

```
Data "table" -> List[Tuple[string, Value]] : SymbolTable
```

or (ii) use a dictionary data structure coming from .NET, such as `ImmutableDictionary`, which was the implementation choice for Casanova. In both cases, the behaviour of the language implemented in Metacasanova will be that of a dynamic language, because whenever the value of a variable or class field must be read, the evaluation rule must look up the symbol table at run time to retrieve the value, whose complexity will be $O(n)$ with the list implementation and $O(\log n)$ with the dictionary implementation. This issue is caused by the fact that, in the current state of Metacasanova, the meta-type system is unaware of the type system of the language that is being implemented in the meta-compiler. This is not a problem limited to Meta-

casanova but to all meta-compilers having a meta-type system that does not allow embedding of the host language type system. In the next section we propose an extension to Metacasanova to overcome this problem by embedding the type system of the implemented language in the meta-type system of Metacasanova and inlining the code to access the appropriate variable at compile time.

Chapter 5

Metacasanova optimization

In Section ?? and Section ?? we presented the semantics of Metacasanova and we showed how the meta-compiler generates the code necessary to represent the elements of the language and the evaluation of the rules expressed in terms of operational semantics. In Section 4.2.5 we highlighted the problem of performance degradation, due to the additional abstraction layer of the meta-compiler, and identified a possible cause in how the language manages the memory representation. For now, the memory can only be expressed with a dynamic symbol table that must be looked up at run-time in order to retrieve the value of a variable or of a class/record field. In this section we propose an extension to Metacasanova with parametric *Modules* and *Functors* that will allow to inline the access to record fields at compile time and to embed an arbitrary type system into the meta-type system of Metacasanova. Note that in this scope, we use the term functor with the same meaning used in the scope of the language CamL, i.e. a function that takes some types as input and returns a type. In order to provide additional clarity to the explanation, we introduce, in the next section, an example that we use as reference across the whole section.

5.1 Case study

Assume that we want to represent a physical body with a **Position** and a **Velocity** in a 2D space. This can be defined as a data structure containing two fields for its physical properties (the example below is written in F#).

```
type PhysicalBody = {  
    Position      : Vector2  
    Velocity      : Vector2  
}
```

In the current state of the Metacompile, a language that wants to support such a data structure, as stated in Section ??, should define it either with a list of pairs (*field, value*) or with a dictionary from .NET.

```
Data "Record" -> List[Tuple[string, Value]] : Record
```

Accessing the values of the fields requires to iterate through this list (or dictionary) and find the field we want to read, with two evaluation rules such as

```
field = name
-----
getField ((field,value) :: fields) name
  -> value

field <> name
getField fields name -> v
-----
getField ((field,value) :: fields) name -> v
```

This could be done immediately by inlining the *getter* (or *setter*) for that field directly in the program.

In what follows we add a system of modules and functors to Metacasanova, we explain how the meta-compiler generates the code for them, and we show how to use them to improve the performance of the example above.

5.2 Using Modules and Functors in Metacasanova

A module definition in Metacasanova is parametric with respect to types, in the sense that a module definition might contain some type parameters, and can be instantiated by passing the specific types to use. A module can contain the definition of data structures, functions, or functors.

```
Module "Record" : Record {
  Functor "RecordType" : * }
```

The symbol *** reads *kind* and means that the functor might return any type. Indeed the type of a record (or class) in a programming language can be “customized” and depends on its specific definition, thus it is not possible to know it beforehand.

We then define two modules for the *getter* and *setter* of a field of a record. In this example, we use type parameters in the module definitions.

```
Module "Getter" => (name : string) => (r : Record) {
  Functor "GetType" : *
  Func "get" -> (r.RecordType) : GetType }

Module "Setter" => (name : string) => (r : Record) {
  Functor "SetType" : *
  Func "set" -> (r.RecordType) -> SetType : (r.RecordType) }
```

These two modules respectively define a functor to retrieve the type of the record field, and a function to get or set its value. Note that in the function definitions *get* and *set* we are calling the functor of the *Record* module to generate the appropriate type for the signature. This is allowed, since the result of a functor is indeed a type.

A record meta-type (i.e. its representation at meta-language level) is recursively defined as a sequence of pairs (*field*, *type*), whose termination is given by *EmptyField*. We thus define the following functors:

```

Functor "EmptyRecord" : Record
Functor "RecordField" => string => * => Record : Record

```

The first functor defines the end point of a record, which is simply a record without fields. The second functor defines a field as the pair mentioned above followed by other field definitions.

Moreover, we must define two functors that are able to dynamically build the *getter* and *setter* for the field.

```

Functor "GetField" => string => Record : Getter
Functor "SetField" => string => Record : Setter

```

The behaviour of functor is expressed, as for normal functions, through a rule in the meta-program. A rule that evaluates a functor returns an instantiation of a module. Note that, inside a module instantiation, it is possible to define and implement functions other than those in the module definition, i.e. the module instantiation must implement *at least* all the functors and functions of the definition. For instance, the following is the type rule instantiating the module for **EmptyRecord**:

```

-----
EmptyRecord => Record {

  Func "cons" : unit

  -----
  RecordType => unit

  -----
  cons -> ()

}

```

The function **cons** defines a constructor for the record, which, in the case of an empty record, returns nothing. The module instantiation for a record field evaluates as well **RecordType**, and has a different definition and evaluation of the function **cons** (because it is constructed in a different way):

```

-----
RecordField name type r = Record {
  Func "cons" -> type -> r.RecordType : RecordType

  -----
  RecordType => Tuple[type,r.RecordType]

  -----
  cons x xs -> (x,xs)}

```

Note that the return type of **cons** is to be intended as calling **RecordType** of the current module, so as it were

this.RecordType. The getter of a field must be able to lookup the record data structure in search of the field and generate a function to get the value from it. For this reason, the functor instantiates two separate modules, depending on the name of the field that we are currently examining.

```

//Rule 1
name = fieldName
thisRecord := RecordField name type r
-----

```

```

GetField fieldName (RecordField name type r) => Getter name thisRecord {
  GetType => type
  -----
  get (x,xs) -> x}

//Rule 2
name <> fieldName
thisRecord := RecordField name type r
-----
GetField fieldName (RecordField name type r) => Getter name type
  thisRecord{
    Functor "GetAnotherField" : Getter
    -----
    GetAnotherField => GetField fieldName r

    GetAnotherField => g
    -----
    GetType => g.GetType

    GetAnotherField => getter
    getter.get xs -> v
    -----
    get (x,xs) -> v }

```

Listing 5.1: Module instantiations for getters

Analogously, the setter of a field instantiates two separate modules whether the current field is the one we want to set or not.

```

name = lt
thisRecord := RecordField name type r
-----
SetField lt (RecordField name type r) => Setter name thisRecord{
  -----
  SetType => type
  -----
  set (x,xs) v -> (v,xs)}

name <> lt
thisRecord := RecordField name type r
-----
SetField lt (RecordField name type r) => Setter name thisRecord{
  TypeFunc "SetAnotherField" : Setter
  -----
  SetAnotherField => SetField lt r
  -----
  SetType => type

  SetAnotherField => setter
  setter.set xs v -> xs'
  -----
  set (x,xs) v -> (x,xs') }

```

Listing 5.2: Module instantiations for setters

5.3 Functor result inlining

If a premise or a conclusion contains a call to a functor, this call is evaluated at compile time, rather than at runtime. Metacasanova has been extended

with an interpreter which is able to evaluate the result of the functor calls. The behaviour of the interpreter follows the same logic explained when presenting the code generation steps in Section ??, thus here we do not present the details for brevity. When a rule outputs the instantiation of the module, the generated code will contain only rules of the modules which conclusion contains a function (i.e. functions that output values, not functors). In this way the generated code will contain a different version of those functions depending on the instantiation parameters of the module.

We now show how to use the implementation of the records given in Section 5.2 for the physical body presented as a case study. The definition of the record type for the physical body is done through a functor

```

Functor "PhysicalBodyType" : Record

EmptyRecord => empty
RecordField "Velocity" Vector2 empty => velocity
RecordField "Position" Vector2 velocity => body
-----
PhysicalBodyType => body

```

This rule is evaluated at compile time by the interpreter that generates one module for each field of the **PhysicalBody**, containing the constructor. For example, for the field **Velocity** the interpreter will generate¹

```

Func "cons" -> Vector2 -> unit : Tuple[Vector2,unit]

-----
cons x xs -> (x,xs)

```

This because the functor will call the evaluation rule for **RecordField** with the argument (**Recordfield** "Velocity" Vector2 (**EmptyRecord**)). This rule generates the function **cons** by evaluating the result of the functors **EmptyRecord.RecordType** and **RecordField.RecordType**, which respectively produce **unit** and **Tuple[Vector2,unit]**.

Instantiating a physical body will just require to build a function that returns the type of the physical body, which is obtained by calling the functor **PhysicalBodyType**.

```

Func "PhysicalBody" : PhysicalBodyType.RecordType

-----
PhysicalBody -> PhysicalBodyType.cons((Vector2.Zero,(Vector2.Zero,())))

```

Defining the setter and getter of a field, requires to use the functor **GetField** to generate the appropriate getter function. After the module has been correctly generated, we can use the getter for the field. For example, in order to get the position field, we use the following function.

```

Func "getPos" -> PhysicalBodyType : Vector2

GetField "Position" PhysicalBodyType => getter
getter.get PhysicalBody -> p
-----
getPos -> p

```

¹Note that here we give a high-level representation of the generated rules that are actually directly generated as C# code.

The result of the premise `GetField` will be evaluated at compile time through the code in Listing 5.1 and will instantiate a module containing the following function definition and rule.

```
Func "get" -> Tuple[Vector2,Tuple[Vector2,unit]] : Vector2
-----
get (x,xs) -> x
```

Note that the second premise of `getPos` will immediately call the `get` generated in this step. The case of `setPos` is analogous except the setter takes an additional argument.

Reading `Velocity` analogously uses a functor call to generate a getter:

```
Func "getVel" -> PhysicalBodyType : Vector2

GetField "Velocity" PhysicalBodyType => getter
getter.get PhysicalBody -> p
-----
getVel -> p
```

This time the functor will generate two different functions in two separate modules. The first time the record is processed, **Rule 2** in Listing 5.1 will be activated (because the first field in the Record is `Position`). This rule will instantiate an additional module when evaluating the functor call in its premise, which in turn is able to get the `Velocity` field. The rule for `get` in the first module will contain in its premise a call to `get` of the second module.

```
//Code for module1
Func "get" -> Tuple[Vector2,Tuple[Vector2,unit]] : Vector2

module2.get xs -> v
-----
get (x,xs) -> v

//Code for module2 generated by evaluating the functor in the premise of
Rule 2
Func "get" -> Tuple[Vector2,unit] : Vector2
-----
get (x,xs) -> x
```

We want to point out that this optimization has been presented on the specific case of records, but can be generalized for any situations where you would use a symbol table. Indeed any symbol table can be expressed with the representation above as a sequence of pair where the first item is the value of the current variable, and the second item is the continuation of the symbol table.

5.4 Functor interpreter

Here describe how the inlining process is implemented in the meta-compiler with the type function interpreter

5.5 C-- optimization

Show how to optimize the current C-- implementation by using functors to populate the symbol table.

5.6 Casanova 2.5 optimization

- Entity definition with functors.
- Entity traversal with functors (?)
- Rule update definition with functors.
- Evaluation.

Show how to use functors to define an entity as a Record and how to inline the getter and setter of fields in rules.

5.7 Evaluation

An extensive evaluation of Casanova implemented in Metacasanova, which we omit for brevity, can be found in [19]. The implementation of Casanova operational semantics in Metacasanova is almost 5 times shorter than the corresponding F# implementation in the hard-coded compiler. In addition to Casanova, we have implemented a subset of the C language called C--. This language supports `if-then-else`, `while-loop`, and `for` statements, as well as local scoping of variables. The total length of the language definition in Metacasanova is 353 lines of code. The corresponding C# code to implement the operational semantics of the language is 3123 lines, thus the code reduction with Metacasanova is roughly 8.84 times. For comparison, in Table 5.2 it is possible to see the code length to implement three different statements, both in Metacasanova and C#. We tested C-- against Python by computing the average running time to compute the factorial of a number. C-- results to be 50 times slower than Python. This result is worse than what we obtained with Casanova, because in order to emulate the interruptible rule mechanism of Casanova in Python you must rely on coroutines that are slower than a program containing simple statements. Moreover, we tested the performance improvement of the optimization using Functors to represent records against the standard one using dynamic symbol tables. The test was run using records with a number of fields ranging from 1 to 10 and updating from 10000 to 1000000 instances of such records. In Table 5.1, we can see that the optimization using Functors leads to a performance increase on average of about 11 times, with peaks of 30 times. The gain increases with the number of fields, thus Functors are particularly effective for records with high number of fields. Figure 5.1 shows a chart of the overall performance of the two techniques (the data points are taken from Table 5.1). The horizontal axis contains the amount of fields per record, while the vertical axis contains the number of records that are being updated. We can see that the performance of the dynamic table degrades considerably when increasing the number of fields, and that the higher the amount of records is, the steeper the curve is. On the other hand, the performance of the implementation with Functors is almost constant, regardless of the amount of fields or records that are being updated. Moreover, note that the




Figure 5.1: Execution time of the different memory models

performance of the dynamic table is improved by the fact that we are using a dictionary implemented in .NET, which can access the entries in $O(\log n)$. If the symbol table were represented as a meta-data structure in the language the performance would be even worse, since it would have to be encoded as a list of pairs with the field name and its value, and its manipulation would be affected by the evaluation rules that should implement this behaviour. Furthermore, the dynamic lookup should be done also to ensure that the types of the record fields are used consistently (for example to prevent that a record is constructed with incompatible values for its fields), while using the functors in Metacasanova embeds the type system of the language in the meta-type system, whose type safety is checked at compile-time rather than at runtime, and this contributes to further increase the performance.

Table 5.1: Running time with the functor optimization and the dynamic table with 10000, 100000, and 1000000 records.

FIELDS	Functors (ms)	Dynamic Table (ms)	Gain
1	1.00E-05	5.00E-06	0.50
2	9.00E-06	1.30E-05	1.44
3	9.00E-06	2.70E-05	3.00
4	9.00E-06	4.50E-05	5.00
5	9.00E-06	7.00E-05	7.78
6	9.00E-06	9.90E-05	11.00
7	9.00E-06	1.33E-04	14.78
8	9.00E-06	1.75E-04	19.44
9	9.00E-06	2.20E-04	24.44
10	9.00E-06	2.70E-04	30.00
Average gain			11.74

FIELDS	Functors (ms)	Dynamic Table (ms)	Gain
1	9.60E-05	6.30E-05	0.66
2	9.40E-05	1.59E-04	1.69
3	9.50E-05	3.04E-04	3.20
4	9.60E-05	5.03E-04	5.24
5	9.60E-05	7.52E-04	7.83
6	9.60E-05	1.05E-03	10.95
7	9.70E-05	1.41E-03	14.57
8	9.80E-05	1.82E-03	18.59
9	9.90E-05	2.29E-03	23.17
10	1.00E-04	2.81E-03	28.05
Average gain			11.39

FIELDS	Functors (ms)	Dynamic Table (ms)	Gain
1	9.47E-04	7.29E-04	0.77
2	9.51E-04	1.78E-03	1.87
3	9.50E-04	3.33E-03	3.51
4	9.60E-04	5.43E-03	5.66
5	9.65E-04	8.03E-03	8.32
6	9.71E-04	1.11E-02	11.44
7	9.75E-04	1.47E-02	15.12
8	9.82E-04	1.89E-02	19.28
9	9.92E-04	2.37E-02	23.86
10	1.00E-03	2.87E-02	28.62
Average gain			11.84

Table 5.2: Code length implementation of C-- and run-time performance

Statement	Metacasanova	C#
if-then-else	4	103
while	7	73
For	11	81

C--	Python
1.26ms	$2.36 \cdot 10^{-2}$ ms

Chapter 6

Networking primitives in Casanova 2

In this section we introduce the basic concepts of the implementation of multiplayer game development for Casanova 2. This implementation aims to relieve the programmer of the complexity of hard-coding the network implementation for an online game, while preserving encapsulation in code. We show that code analysis is required to generate the appropriate network primitives to send and receive data. Finally, we present a simple multiplayer game to show a concrete example.

6.1 Introduction

Adding multi-player support to games is a highly desirable feature. By letting players interact with each other, new forms of gameplay, cooperation, and competition emerge without requiring any additional design of game mechanics [27]. This allows a game to remain fresh and playable, even after the single player content has been exhausted. For example, consider any modern AAA (AAA refers to games with the highest development budgets[44]) game such as *Halo 4*. After months since its initial release, most players have exhausted the single player, narrative-driven campaign. Nevertheless the game remains heavily in use thanks to multiplayer modes, which in effect extended the life of the game significantly. This phenomenon is even more evident in games such as *World of Warcraft* or *EVE*, where multiplayer is the only modality of play and there is no single-player experience.

Challenges Multi-player support in games is a very expensive piece of software to build. Multiplayer games are under strong pressure to have very good *performance* [16]. Performance is both expressed in terms of CPU time and in bandwidth used. Also, games need to be very *robust* with respect to transmission delays, packets lost, or even clients disconnected. To make matters worse, players often behave erratically. It is widespread practice among players to leave a competitive game as soon as their defeat is apparent (a phenomenon so common to even have its own name: “rage

quitting” [30]), or to try to abuse the game and its technical flaws to gain advantages or to disrupt the experience of others.

Networking code reuse is quite low across titles and projects. This comes from the fact that the requirements of every game vary significantly: from turn-based games that only need to synchronize the game world every few seconds, and where latency is not a big issue, to first-person-shooter games where prediction mechanisms are needed to ensure the smooth movement of synchronized entities, to real-time strategy games where thousands of units on the screen all need to be synchronized across game instances [40]. In short, previous effort is substantially inaccessible for new titles.

Encapsulation suffers from this ad-hoc nature of the implementation of the networking layer in multiplayer games. Indeed managing the information about game updates over a network requires each game entity to interface the game logic code with network connection and socket objects, data transmission method calls such as send and receive, and support data structures to manage traffic and track the status of common protocols. This happens because each game entity must provide the following functionality in order to work in a multiplayer game:

- Update the logic in the fashion of a singleplayer counterpart.
- Choose what data is necessary to send over the network and create the message containing this information.
- Choose what data can be lost and what data must always be received by the other clients.
- Periodically check if incoming messages contain information that needs to be read and to perform specific updates.

Combining these requirements together within the same entity breaks encapsulation because now the logic of the entity and lots of spurious details only relevant to the networking implementation are mixed together, resulting in a highly noisy program. Maintenance then becomes very hard, as every change in the game logic must also be reflected in the networking implementation.

Existing approaches Networking in games is usually built with either very low-level or very high-level mechanisms. Very low-level mechanisms are based on manually sending streams of bytes and serializing only the essential bits of the game world, usually incrementally, on unreliable channels (UDP). This coding process is highly expensive because building by hand such a low-level protocol is difficult to get right, and debugging subtle protocol mismatches, transmission errors, etc. will take lots of development resources. Low-level mechanisms must also be very robust, making the task even harder.

High-level protocols such as RDP, reflection-based serialization, frameworks (such as Pastry, netty.io), etc. can also be used. These methods greatly simplify networking code, but are rarely used in complex games and scenarios. The requirements of performance mean that many high-level protocols or mechanisms are insufficient, either because they are too slow computationally (especially when they rely on reflection or events) or because they transmit too much data across the network.

6.2 Motivation

To avoid the problems of both existing approaches, we propose a middle ground. We observe that networking fundamental abstractions upon which the actual code and protocols are built do not vary substantially between games, even though the code that needs to be written to implement them does. The similarity comes from the fact that the ways to serialize, synchronize, and predict the behaviour of entities are relatively standard and described according to a limited series of general ideas. The difference, on the other hand, comes from the fact that low-level protocols need to be adapted to the specific structure of the game world and the data structures that make it up. Until now, common primitives have not been syntactically and semantically captured inside existing domain-specific languages for game development [12]. Using the right level of abstraction, these general patterns of networking can be captured, while leaving full customization power in the hand of the developer (to apply such primitives to any kind of game).

6.3 Related work

In the following we discuss some existing networking tools used in game development and we highlight some issues that arise from their use.

The Real time framework (RTF) RTF [26] is a middleware built for C++ to relieve the programmer from dealing with data compression. It is more flexible than solutions based on game engines or hand-made implementations, since it automates the process of data transmission. Moreover, it supports distributed server management. Unfortunately, this solution has several flaws:

- All entities must inherit from the class `Local` and the semantics of the position is pre-determined, often clashing with rendering or physics.
- Platform independence requires that the programmer uses RTF primitive types.
- Data transmission automation requires that all game entities inherit the class `Serializable`.
- Being a middleware, RTF is not aware of what games are going to use it for (every game comes with different data structures). Thus, the developer is tasked to include in his code also logic to update the RTF layer, in order to keep the game updated over the network.

Network scripting language (NSL) NSL [38] provides a language extension based on a send-receive mechanism. Moreover it provides a built-in client side prediction (a feature missing in existing highly concurrent and distributed languages such as Stackless Python [42] and Erlang [9]), which is periodically corrected by the server.

Unreal Engine/Unity Engine Unreal Engine [2] and Unity Engine [1] are commercial game engines supporting networking. Both Unity and Unreal Engine use a client-server approach. In Unreal Engine, the server contains the “true” game state, and the clients contain a “dirty” copy, which is validated periodically. It is possible to define entities (actors in Unreal Engine jargon) that are replicated on the clients. Whenever a replicated actor changes on the server, this change is also reflected on the clients. Additional customization can be achieved through Remote procedure calls (RPCs) of three kinds.

- The function is called on the server and executed on the client. This is useful for game elements that do not affect gameplay, such as creating a particle effect when a weapon is fired.
- The function is called on the client and executed on the server. This is useful for events that affect the other clients and should be validated by the server.
- The function is executed in multi-cast, meaning that the server calls the function and that it is executed on both the server and all the clients.

The Unity Engine uses a similar approach based on networking components, synchronized at every frame, and RPC’s to define custom synchronization events.

Unfortunately, customization comes at the cost of the level of detail that developers must face. Using RPC’s require a deep knowledge of the engine and writing lots of code, as discussed in Section ??.

In this section we introduce a small example that addresses the requirements of designing a multiplayer game. We then present an architecture that aims to fulfil these requirements.

6.4 The master/slave network architecture

We chose to implement the networking layer in Casanova 2 by using a peer-to-peer architecture for the following reasons:

- Server-client architectures are more reliable but suitable only for specific genres of games (mostly Shooter games), while other genres, such as Real-time strategy games or Online Role Playing Games, use P2P architectures.
- We do not have to write a separate logic for an authoritative game server, which has to validate the actions of clients.

Casanova will provide a generic tracking server, which is run separately from the main program. The tracking server is a thin service that connects players participating in a single game, and helps with forwarding the network traffic through NATs (Network Address Translation).

Each client maintains a local copy of the `world` entity and has direct control over a single portion of it. Instances belonging to such a portion are seen as *master* by this player, who is always allowed to directly change


the state of the master instances without having to validate this state change by synchronizing with other players through the network.

Each client also maintains a portion of the world that is not directly under his control. Instances belonging to such as portion are seen as *slave* by this player, who is only allowed to *predict* the local state of the instances and, whenever he receives an update from their masters, must correct this prediction according to the data contained in the received messages. The slave part of the world is thus maintained passively by the client: the only active part is predicting the evolution of the entity state and correcting it whenever he receives an update by its master.

For this purpose, we extend the syntax of Casanova rules by allowing them to be marked with the modifiers **master** and **slave**. These rules are executed respectively on master and slave entities. Note that it is still possible not to mark a rule with these modifiers, which means that the rule is always executed independently of the fact that the entity is either master or slave on that particular client. We also allow to mark a rule as **connecting** and **connected**. These rules are triggered only once respectively when a new client connects and when the clients detect a new connection.

Casanova also provides primitives to send (reliably or unreliably) and receive data. A schematic representation of this architecture can be seen in Figure 6.2.

Figure 6.1: Representation of the game world in a networking scenario



Note the aim of this architecture is to provide language-level primitives to describe the networking logic. This means that the compiler will be able to generate code compatible with the low-level network libraries that provide transmission functions over the network channel without having to change Casanova code in the program. In our implementation, we chose the .NET library *Lidgren*, which is widely used also in commercial game engines such as Unity3D and MonoGame, but nothing prevents the compiler to be expanded in order to target other similar libraries for other languages, such as jgroups [10].




Figure 6.2: master/slave architecture

6.5 Case study

Let us consider a simple shooter game where each player controls a space ship. Players can move forward, backward, and rotate the ship to change direction. Moreover, they can use the ship lasers to shoot other players. If a laser hits an enemy ship, we increase the player's score. Designing such a game requires to address the following issues, depicted by the schematic representation in Figure 6.1:

1. Each player must maintain a local version of the game state (world). In order to avoid to flood the network with messages, all the copies are not fully synchronized at each frame, thus they are slightly different and each client knows the latest version of only part of the copy.
2. A player **connecting** to an existing game must be able to receive the latest update of the game state and send the new ship he will control to existing players in the game.
3. A player already **connected** to the game must detect a new connection and send his master portion of the game state.
4. Each player must be able to control only one ship at a time. This means that the part of the game logic that processes the input and modifies the spatial data of the ship (position and rotation) should only be executed on the ship controlled by the player and not on the local copies of other players' ships. This means that each player sees as **master** only one ship instance.
5. Each player must send the updated state of the ship he controls to the other players after executing the local update. To achieve better performance over the network, the data is not sent at every update, but with a lower frequency.

6. Each player must receive the updated state of `slave` ships controlled by other players. In this phase, we must take into account that, as explained above, not every update is sent, so the player should “predict” what will happen during the game frames in which he does not receive an update.

6.6 Implementation

Each of the scenarios described above requires specific language extensions. These extensions identify connection, ownership (master/slave), and various send and receive primitives. In this section, we introduce each primitive by using a multiplayer game example ¹. We now give an implementation of the shooter game presented above, using the extended version of Casanova 2 with network primitives.

The `world` contains a list of ships controlled by each player.

```
world Shooter = {  
  Ships : [Ship]  
  ...  
}
```

Each `Ship` contains a position, a rotation, a collection of shot projectiles, and the score.

```
entity Ship = {  
  Position : Vector2  
  Rotation : float32  
  Projectiles : [Projectile]  
  Score : int  
  ...  
}
```

Each `Projectile` contains its position and velocity.

```
entity Projectile = {  
  Position : Vector2  
  Velocity : Vector2  
  ...  
}
```

Connection

When a player connects, we must consider two different situations: (i) a player is already in the game and must send the current game state to the connecting players, and (ii) the player who is connecting needs to send the ship he will instantiate and control (its initial state). Both the players in the game and the connecting one must receive the game states that are sent. For this purpose we introduce two additional modifiers, `connecting` and `connected`, that can be added to rule declarations to mark their role in the multiplayer logic.

¹The game source code and executable can be found at <https://github.com/vs-team/casanova-mk2/wiki/Networking-extension>

Connecting A rule marked with **connecting** is executed once when a player joins the game for the first time. In our example, the player should send his initial state (the created ship) to the other players. We use the primitive **send_reliable** because we must be sure that eventually all players will be notified of the ship creation.

```
world Shooter = {
...
rule connecting Ships =
yield send_reliable Ships
}
```

Connected A rule marked with **connected** is run whenever a new player joins the game by all existing players. When this occurs, each player sends its ship. The system will take care to send only the ship controlled locally by the player itself for each player. The rule will use the **send_reliable** primitive for the same reason explained in the previous point.

```
world Shooter = {
...
rule connected Ships =
yield send_reliable Ships
}
```

Note that even if the code is the same, the semantics of the two rules are different. The first one is executed by the player joining the game, who locally instantiates its **Ship** and must send its list of **Ships** (containing only the local instance) to the other players. The second one is executed by all existing players who must share with the joining player the list of existing ships.

Master updates

As explained above, each client manages a series of local game objects (called *master objects*) that are under its direct control. The other clients read passively any update done on those instances and update their remote copy (*slave objects*) accordingly. We mark rules affecting the behaviour of master objects as **master**. In our example, the following situations are run as master: (i) synchronizing the ships among players, (ii) updating the ship and projectiles spatial data, and (iii) creating and destroying projectiles.

1. Each player is tasked to maintain the list of **Ships** in the world. This requires to receive the updated list from other players and to store the new value in a master rule. Indeed the world is a special case of an entity that is shared among players, and not directly owned by somebody. Each ship contained in that list and received from other players will be treated appropriately as slaves, while the only one owned by the current player will be under his direct control. In this rule we use **let!**, which is an operator that waits until the argument expression returns a result and then binds it to the variable. The symbol **@** stands for list concatenation. The rule uses **receive_many**, which receives and collects the list of sent ships by the other players.

```

world Shooter = {
  ...
  rule master Ships =
    let! ships = receive_many()
    yield Ships @ ships
}

```

2. The master version of the ship update reads the input of the player and moves (or rotates) the ship if the appropriate key is pressed. Note that this part must be executed only on a master object, because we want to allow each player to control only the ship he owns and instantiates at the beginning of the game. Below we show just the rule to move forward; the other movement and rotation rules are analogous. We use an *unreliable send* because it is acceptable to lose an update of the position during a certain frame: shortly after, there will be a new update.

```

entity Ship = {
  ...
  rule master Position =
    wait world.Input.IsKeyDown(Keys.W)
    let vp = new Vector2(Math.Cos(Rotation),
      Math.Sin(Rotation)) * 300.0f
    let p = Position + vp * dt
    yield send p
}

```

We do the same for projectiles, except the projectile position is continuously updated and synchronized over the network without having to wait that a key is pressed.

3. Creating a new projectile happens when the player shoots. A ship keeps track of the projectiles it has shot so far, and adds a new one to the list of the existing projectiles. The updated list is sent to all players with the new instance of the projectile (which is added as a new head of the list with the operator `::`). Here it is better to precise the semantics of the `yield` in conjunction with the use of networking primitives. A `yield` requires that the written value is type-compatible with the domain of the rule. Thus, when used with a `send` primitive, we must pass as argument a list. The system will ensure, for performance reasons, that the generated code only sends the new items added to the list. This semantics is defined as such for two main reasons: (i) when sending the new projectiles we must also update the list in local (and given the immutability of Casanova we must replace the existing one), and (ii) because in this way the programmer can focus on the logic of the game as if it were a single-player game without worrying of network-specific details. Note that the last `wait` forces the player to release the key before shooting again (semi-automatic fire). Removing that check would spawn multiple projectiles consecutively, which is not a wanted behaviour.

```

entity Ship = {
  ...
  rule master Projectiles =

```

```

wait world.Input.IsKeyDown(Keys.Space)
let vp = new Vector2(Math.Cos(Rotation),
Math.Sin(Rotation)) * 500.Of
let proj = new Projectile(Position, vp) :: Projectiles
yield send_reliable proj
wait not world.Input.IsKeyDown(Keys.Space)
}

```

Filtering the colliding projectiles and updating the score is run as a master rule. The rule computes the set difference between the ship projectiles and the colliding projectiles and updates the list of projectiles, sending them through the network as well. Even in this case, the network layer sends only the information about the projectiles to remove. Note that the score is managed by each player locally, as it does not require to be synchronized (we do not print the other players' scores. Doing so would indeed require to also send the score).

```

entity Ship = {
...
rule master Projectiles, Score =
let collidingProjs =
[for p in Projectiles do
let ships =
[for s in Ships do
where
s <> this and
Vector2.Distance(p.Position,s.Position) < 100.Of
select s]
where ships.Count > 0
select p]
let newProjectiles = Projectiles - collidingProjs
yield send_reliable newProjectiles,
Score + collidingProjs.Count
}

```

Managing remote instances

The game objects that were not instantiated by a client, but received from another client, are *slave objects* and must be synchronized differently than master objects. For this purpose, a rule can be marked as **slave**. In our example, we use slave rules in the following situations: (i) synchronizing other players' ships and projectiles spatial data, and (ii) projectiles instantiated by other players.

1. Every remote projectile and ship is synchronized locally by a rule, which tries to **receive** a message containing updated spatial data. Below we provide the code to update the position of the ship; the synchronization of other spatial data is analogous.

```

entity Ship = {
...
rule slave Position = yield receive()
}

```

2. When a projectile is instantiated remotely, we have to receive it and add it to the list of projectiles. We use **receive_many** because the new projectiles are added to a list. This case also supports the situation where a ship could shoot multiple projectiles at the same time.

```
entity Ship = {  
  ...  
  rule slave Projectiles =  
  let! projs = receive_many()  
  yield projs @ Projectiles  
}
```

In this scenario is important to discuss the atomicity of these transmissions: in the context of network games, reliability is often sacrificed for better network performance, so most of the data transmissions are unreliable (like in the case of the ship position). This means that we have no guarantee that the message will be received. Several issues can arise from this situation: for example, if a player fails to receive the position of the ship, then it might miss a collision with a projectile. This is a well-known issue in several shooter games and out-of-sync errors might happen during a multiplayer game. However, ensuring that all the data transmissions are reliable might affect network performance to the point that the game would become unplayable because of the network overload.

Casanova 2 allows the programmer to decide whether the transmission should be reliable or not and experiment with the effect of a reliable transmission versus an unreliable one that does not overload the network. For example, the updated list of projectiles, after a collision, is always sent in a reliable way. This is acceptable because collisions are not so frequent. This is not true for the ship position, since movements are very frequent and mostly happen at every frame, thus it is something that should not be sent reliably at every frame.

Furthermore, we want to focus the attention on the implicit relationship between this networking architecture and the encapsulation: as shown for instance in the examples where the ship shoots a projectile, we ensure encapsulation by keeping a semantics that filters completely the details about networking. The programmer only worries about the logic of adding a new projectile, while the details of the network transmission are hidden. A hand-made implementation is usually prone to break this separation of concerns because the transmission logic is tightly coupled within the game logic itself.

6.7 Networking in Metacasanova

Chapter 7

Discussion and conclusion

Here it would be wise to discuss about the fact that you still need to define a program for a language implemented in Metacasanova in term of syntax of Metacasanova itself. Remark that it is trivial to solve this problem by writing a parser, for example in Yacc, that maps the proper syntax of the language into the syntax of Metacasanova, but we do not implement it because it has no scientific value.

Appendix A

List operations with templates

Bibliography

- [1] Unity Game Engine. <http://unity3d.com>.
- [2] Unreal Technology. <https://www.unrealengine.com/>.
- [3] Mohamed Abbadi. *Casanova 2, A domain specific language for general game development*. PhD thesis, Università Ca' Foscari, Tilburg University, 2017.
- [4] Mohamed Abbadi, Francesco Di Giacomo, Agostino Cortesi, Pieter Spronck, Giulia Costantini, and Giuseppe Maggiore. Casanova: A simple, high-performance language for game development. In Stefan Göbel, Minhua Ma, Jannicke Baalsrud Hauge, Manuel Fradinho Oliveira, Josef Wiemeyer, and Viktor Wendel, editors, *Serious Games*, volume 9090 of *Lecture Notes in Computer Science*, pages 123–134. Springer International Publishing, 2015.
- [5] Mohamed Abbadi, Francesco Di Giacomo, Agostino Cortesi, Pieter Spronck, Giulia Costantini, and Giuseppe Maggiore. Casanova: a simple, high-performance language for game development. In *Joint International Conference on Serious Games*, pages 123–134. Springer, 2015.
- [6] Alfred V Aho, Ravi Sethi, and Jeffrey D Ullman. Compilers, principles, techniques. *Addison wesley*, 7(8):9, 1986.
- [7] Alfred V Aho, Ravi Sethi, and Jeffrey D Ullman. *Compilers: principles, techniques, and tools*, volume 2. Addison-wesley Reading, 2007.
- [8] Andrew W. Appel. *Modern compiler implementation in Java*. Cambridge University Press, 2002.
- [9] Joe Armstrong, Robert Virding, Claes Wikström, and Mike Williams. *Concurrent programming in ERLANG*. Prentice Hall, 1993.
- [10] B. Ban. JGroups - A Toolkit for Reliable Multicast Communication. <http://www.jgroups.org/index.html>, 2002.
- [11] Henk Barendregt, Wil Dekkers, and Richard Statman. *Lambda calculus with types*. Cambridge University Press, 2013.

- [12] S. Bhatti, E. Brady, K. Hammond, and J. McKinna. Domain specific languages (DSLs) for network protocols. In *International Workshop on Next Generation Network Architecture (NGNA 2009)*, 2009.
- [13] Martin Bravenboer, Karl Trygve Kalleberg, Rob Vermaas, and Eelco Visser. Stratego/xt 0.17. a language and toolset for program transformation. *Science of computer programming*, 72(1):52–70, 2008.
- [14] Luca Cardelli. Type systems. *ACM Computing Surveys*, 28(1):263–264, 1996.
- [15] Alonzo Church. A formulation of the simple theory of types. *The journal of symbolic logic*, 5(2):56–68, 1940.
- [16] M. Claypool and K. Claypool. Latency and player actions in online games. *Communications of the ACM*, 49(11):40–45, 2006.
- [17] Krzysztof Czarnecki and Ulrich W Eisenecker. *Generative programming: methods, tools, and applications*, volume 16. Addison Wesley Reading, 2000.
- [18] Nils Anders Danielsson and Ulf Norell. Parsing mixfix operators. In *IFL*, pages 80–99. Springer, 2008.
- [19] Francesco Di Giacomo, Mohamed Abbadi, Agostino Cortesi, Pieter Spronck, and Giuseppe” Maggiore. *Intelligent Technologies for Interactive Entertainment: 8th International Conference, INTETAIN 2016, Utrecht, The Netherlands, June 28–30, 2016, Revised Selected Papers*, chapter Building Game Scripting DSL’s with the Metacasanova Metacompiler, pages 231–242. Springer International Publishing, Cham, 2017.
- [20] Atze Dijkstra, Jeroen Fokker, and S Doaitse Swierstra. The architecture of the utrecht haskell compiler. In *Proceedings of the 2nd ACM SIGPLAN symposium on Haskell*, pages 93–104. ACM, 2009.
- [21] E. Dijkstra. Algol-60 translation. Technical report, Mathematisch Centrum - Amsterdam, November 1961.
- [22] Matthias Felleisen and Robert Hieb. The revised report on the syntactic theories of sequential control and state. *Theoretical Computer Science*, 103(2):235 – 271, 1992.
- [23] Firaxis Games. Civilization iv scripting api reference. http://wiki.massgate.net/Our_Python_files_and_Event_Structure, October 2008.
- [24] Plotkin G.D. A structural approach to operational semantics. Technical report, Computer science department, Aarhus University, 1981.
- [25] Francesco Di Giacomo. Casanova 2.5 source code. <https://github.com/vs-team/metacompiler/tree/master/Sources/Content/Content/CNV3>, 2016.

- [26] F. Glinka, A. Ploß, J. Müller-Iden, and S. Gorlatch. RTF: a real-time framework for developing scalable multiplayer online games. In *Proceedings of the 6th ACM SIGCOMM workshop on Network and system support for games*, pages 81–86. ACM, 2007.
- [27] C. Granberg. *David Perry on game design: a brainstorming toolbox*. Cengage Learning, 2014.
- [28] Cordelia Hall, Kevin Hammond, Will Partain, Simon L Peyton Jones, and Philip Wadler. The glasgow haskell compiler: a retrospective. In *Functional Programming, Glasgow 1992*, pages 62–71. Springer, 1993.
- [29] Graham Hutton and Erik Meijer. Monadic parsing in haskell. *Journal of functional programming*, 8(4):437–444, 1998.
- [30] E. Kaiser and W. Feng. PlayerRating: a reputation system for multiplayer online games. In *Proceedings of the 8th Annual Workshop on Network and Systems Support for Games*, page 8. IEEE Press, 2009.
- [31] Jan Willem Klop et al. Term rewriting systems. *Handbook of logic in computer science*, 2:1–116, 1992.
- [32] Massive Entertainment. World in conflict script reference. <http://civ4bug.sourceforge.net/PythonAPI/>, September 2007.
- [33] Robert McNaughton and Hisao Yamada. Regular expressions and state graphs for automata. *IRE transactions on Electronic Computers*, (1):39–47, 1960.
- [34] John C Mitchell. *Concepts in programming languages*. Cambridge University Press, 2003.
- [35] Mikael Pettersson. A compiler for natural semantics. In *Compiler Construction*, pages 177–191. Springer, 1996.
- [36] Benjamin C Pierce. *Types and programming languages*. MIT press, 2002.
- [37] Willard Van Orman Quine, Patricia S Churchland, and Dagfinn Føllesdal. *Word and object*. MIT press, 2013.
- [38] G. Russell, A.F. Donaldson, and P. Sheppard. Tackling online game development problems with a novel network scripting language. In *Proceedings of the 7th ACM SIGCOMM Workshop on Network and System Support for Games*, pages 85–90. ACM, 2008.
- [39] Tim Sheard and Simon Peyton Jones. Template meta-programming for haskell. In *Proceedings of the 2002 ACM SIGPLAN workshop on Haskell*, pages 1–16. ACM, 2002.
- [40] J. Smed, T. Kaukoranta, and H. Hakonen. Aspects of networking in multiplayer computer games. *The Electronic Library*, 20(2):87–97, 2002.

- [41] Harald Søndergaard and Peter Sestoft. Referential transparency, definiteness and unfoldability. *Acta Informatica*, 27(6):505–517, 1990.
- [42] C. Tismer. Stackless Python. <http://www.stackless.com/>.
- [43] Philip Wadler. Monads for functional programming. In *International School on Advanced Functional Programming*, pages 24–52. Springer, 1995.
- [44] M.J.P. Wolf. *The video game explosion: a history from PONG to Playstation and beyond*. ABC-CLIO, 2008.